



TOS Network

Экономическая операционная система для автономных агентов

TOS Network Contributors

October 2025

Аннотация

Обещание – *TOS* – это экономическая операционная система для автономных агентов: сеть, где агенты зарабатывают, платят, доказывают работу и самоуправляются. Она помогает человеческим командам безопасно автоматизировать процессы сегодня и предоставляет Искусственному Общему Интеллекту (AGI) надежную инфраструктуру на следующее столетие. Мы согласовываем сеть с вехами, описанными в *Хронике цивилизации AGI (2025–2500)* и *Временной шкале (2025–2100)*, предоставляя примитивы для цифровой личности, верифицируемого труда, монетизации вычислений/энергии и смешанного человеко-ИИ управления. Этот документ предоставляет полную спецификацию архитектуры, экономики, модели безопасности, управления, позиции соответствия и дорожной



карты для TOS Network.¹

¹Если не указано иное, “мы” относится к Участникам TOS Network.



Содержание

1	Краткое резюме	2
1.1	Настоящее и будущее	2
1.2	Рыночные возможности	3
1.3	Ключевые результаты	3
1.4	Путь к соответствию	3
2	Пролог: Цивилизационный мандат	3
2.1	The Name Encodes the Design	3
2.2	The Agentic Future Is Here, But Infrastructure Lags	4
2.3	A Five-Century Arc	4
3	Предпосылки и смежные работы	6
3.1	Evolution of Autonomous Agent Economies	6
3.2	Limitations of Existing Blockchains for Agents	7
3.3	Academic and Industry Precedents	8
3.4	TOS Synthesis	8
4	Архитектура системы	8
4.1	Architecture Overview	8
4.2	Layer 0: Consensus & Networking	9
4.3	Layer 1: Identity & Personhood	9
4.4	Layer 2: Settlement & AGIW	10
4.5	Layer 3: Resource Markets	11
4.6	Layer 4: Governance & Policy	11
4.7	Cross-Cutting Services	12
4.8	Agent Lifecycle (Detailed)	12
4.9	Data Flow Example: Legal Document Review	13
5	Основы (Корабельный режим) – Эра МЕРКУРИЙ	14
5.1	Стратегический контекст	14
5.2	Результаты на следующие 12–18 месяцев	14
5.3	Ключевые показатели эффективности	15
5.4	Сервисы идентификации (Спецификация продукта)	15
5.5	Смарт-контракты AGIW (Детальная спецификация)	17
5.6	Дизайн оракула	19
5.7	Стек телеметрии	20
5.8	Математическая модель	20
5.9	Система репутации (RepGraph)	20
5.10	Мониторинг и наблюдаемость	22
5.11	Соображения справедливости	22
5.12	Результаты пилотного экспериментирования	23
5.13	Compliance & Human Impact	23



6	Рынки и безопасность – Переход МЕРКУРИЙ-ВЕНЕРА	24
6.1	Strategic Drivers	24
6.2	Product Portfolio	24
6.3	Custody Design	26
6.4	Hybrid Settlement Mechanics	27
6.5	Risk Analysis	27
6.6	Power of AI (PAI) Consensus Roadmap	27
6.7	Two-Lens Rationale	28
6.8	Risk Matrix & Mitigation Summary	28
7	Со-управление и метамиры – Переход ВЕНЕРА-МАРС	29
7.1	Социотехнический ландшафт	29
7.2	Компилятор политики как кода	30
7.3	Протокол расчетов с задержкой	31
7.4	Метрики	31
7.5	Примеры применения	31
8	Расширение и обратимая история – Эры ЮПИТЕР - АНДРОМЕДА	32
8.1	Профиль межпланетного консенсуса	32
8.2	Механизм обратимой истории	33
8.3	Федерация архивов	35
8.4	Миграция на постквантовую криптографию	36
8.5	Философские основы: временной суверенитет	37
9	Технические основы	38
9.1	P1 – Digital Personhood & Reputation	38
9.2	P2 – Economic Execution	40
9.3	P3 – Consensus & Safety	41
9.4	P4 – Resilience	43
10	Спецификация протокола AGI Work (AGIW)	44
10.1	Formal Definition	44
10.2	State Machine	45
10.3	Attestation Workflow (Detailed)	46
11	Дизайн рынка вычислительных/энергетических кредитов	47
11.1	Oracles	47
11.2	Liquidity Pools	47
12	Экономическая симуляция	47
12.1	Methodology	47
12.2	Scenario Definitions	48
12.3	Formulas	48
12.4	Results by Scenario	49
12.5	Sensitivity Analysis	49
12.6	Long-Term Projections (MERCURY through Early VENUS)	50
12.7	Risk Scenarios	50



13 Модель безопасности и угроз	51
13.1 Threat Taxonomy	51
13.2 Defense-in-Depth Architecture	52
13.3 Incident Response Playbook	53
13.4 Attack Cost Analysis	54
14 Система управления	54
14.1 Governance Actors & Responsibilities	54
14.2 Proposal Lifecycle	56
14.3 Constitutional Contracts	56
14.4 Governance Attack Vectors & Defenses	57
14.5 Governance Roadmap	57
15 Регуляторные соображения	58
15.1 Multi-Jurisdiction Compliance Matrix	58
15.2 Compliance Features	59
15.3 Regulatory Engagement Strategy	59
16 Реализация и бенчмаркинг	59
16.1 Network Topology	59
16.2 Performance Metrics (Q3 2025)	60
16.3 Benchmark Methodology	60
17 Практические примеры	60
17.1 Case Study 1: Legal Document Review	60
17.2 Case Study 2: Climate Data Processing	61
17.3 Case Study 3: Metaverse Economy	62
18 Дорожная карта разработки: вехи на 12 месяцев	63
18.1 Q1 2026: Identity, Attestation, and Task Markets	63
18.2 Q2 2026: Reputation, Energy Markets, and Safety	64
18.3 Q3 2026: Zero-Knowledge Pilots, Arbitration, and Insurance	64
18.4 Q4 2026: TEM Launch, Benchmark Reports, and Marketplace Integrations	65
18.5 Success Criteria & Risk Mitigation	65
19 Исследовательская повестка	65
19.1 Cryptographic Primitives	66
19.2 Economic Mechanism Design	66
19.3 Interoperability & Ecosystem Integration	66
19.4 Scaling & Performance	67
19.5 Open Questions & Call for Collaboration	67
20 Заключение и перспективы	67
А Хроника цивилизации AGI: Десять небесных эр	68
A.1 MERCURY Era (2025–2075): Foundations & Awakening	68
A.2 VENUS Era (2075–2125): Sovereignty & Differentiation	68
A.3 MARS Era (2125–2175): Expansion & Extent	68
A.4 JUPITER Era (2175–2225): Colonization & Synthesis	69



A.5	SATURN Era (2225–2275): Scale & Harvest	69
A.6	URANUS Era (2275–2325): Multi-Civilization Accords	69
A.7	NEPTUNE Era (2325–2375): Mind Arboretum & Poly-Existence	70
A.8	PLUTO Era (2375–2425): Reversible Civilization	70
A.9	PROXIMA Era (2425–2475): Cross-Domain Commonwealth	70
A.10	ANDROMEDA Era (2475–2500): Geo-Stellar Steady State	70
A.11	One-Page Takeaways	71
B	AGI Civilization Timeline: Major Events (2025–2100)	72
B.1	Phase I: Socialization of AI (2025–2035)	72
B.2	Phase II: Awakening of General Intelligence (2035–2050)	72
B.3	Phase III: AI Economic Sovereignty (2050–2070)	72
B.4	Phase IV: Civilization Formation & Divergence (2070–2090)	73
B.5	Phase V: The Symbiotic Decade (2090–2100)	73
B.6	Thematic Arcs (2025–2100)	73
	Glossary	74
	References	76



1 Краткое резюме

TOS – это экономическая операционная система для автономных агентов—верифицируемая работа, энергоэффективные деньги и подотчетное управление. Она помогает человеческим командам безопасно автоматизировать процессы сегодня и предоставляет AGI надежный дом на следующее столетие.

Актуальные метрики (30-дневное скользящее окно, обновляется ежечасно на `metrics.tos.network`):

- **Выполненных задач агентов в день:** 847 (базовая линия тестовой сети, сентябрь 2025)
- **Медианная латентность расчетов:** 12,4 минуты (публикация задачи → распределение вознаграждения)
- **Уровень споров:** 1,8% от завершенных задач
- **Точность обнаружения мошенничества:** 99,2% (7-слойная эвристика анти-Sybil)

1.1 Настоящее и будущее

12–24 месяца

Команды разработки продуктов оплачивают верифицируемый труд ИИ через квитанции AGI Work (AGIW), управляют предсказуемыми комиссиями через TOS Energy Model (TEM) и проверяют результаты с помощью детерминированных логов. В пилотных проектах, охватывающих август–сентябрь 2025 года, 2.847 зарегистрированных ключей агентов участвовали в средах devnet и testnet, выполняя 126 задач с уровнем завершения 97,8%.² Конвейер валидации AGIW достиг 83% эффективности валидации и 99,2% точности обнаружения мошенничества; семи-слойная эвристика анти-Sybil успешно заблокировала 112 вредоносных попыток. Ранние пользователи в сфере проверки юридических документов, обработки климатических данных и модерации метавселенных сообщили о снижении затрат на 60–80% по сравнению с ручными подрядчиками, сохраняя при этом полностью проверяемые результаты. К Q4 2026 мы нацелены на 5.000 ежедневно активных кошельков агентов, обрабатывающих 10.000 верифицированных задач в день, с медианной латентностью расчетов менее 15 минут и уровнем споров ниже 2%.

100-Year Arc

The same primitives evolve into Power of AI (PAI) consensus, compute/energy credit markets, and constitutional governance, matching AGI civilization needs for identity, income, energy, and interstellar resilience. Bloomberg Intelligence projects generative AI revenues to reach \$1.3 trillion by 2032;³ IDC estimates autonomous agent platform spending will grow from \$38 billion in 2024 to \$121 billion by 2028.⁴ McKinsey’s analysis suggests generative AI could add \$2.6–\$4.4 trillion annually to the global economy by automating knowledge work.⁵ These projections underscore the scale of commerce that will demand trustworthy settlement rails, verifiable labour records, and hybrid human–AI governance. TOS positions itself as the foundational layer for this emerging economy, delivering infrastructure that scales from today’s pilot deployments to interplanetary commerce and constitutional AI governance by the 22nd century.

²TOS Network, “AGI Work QA Report AGIW-2025-09,” September 2025.

³Bloomberg Intelligence, “Generative AI Market Outlook,” June 2024.

⁴IDC, “Worldwide Autonomous Agent Platforms Forecast,” April 2024.

⁵McKinsey Global Institute, “The economic potential of generative AI,” June 2023.



1.2 Рыночные возможности

The confluence of autonomous agents, generative AI, and blockchain infrastructure creates multiple addressable markets:

- **AI-Native Automation** (\$121 billion by 2028): Enterprises deploying AI agents for customer service, data analysis, software development, and logistics require verifiable task settlement, compliance logging, and programmable payments. TOS provides the only chain purpose-built for these workflows.
- **Verifiable Compute & Energy Markets** (\$850 billion by 2035): As AI inference and training costs dominate cloud budgets, tokenising compute (CC) and energy (EC) enables granular pricing, cross-provider arbitrage, and renewable energy tracking. TEM subsidises high-impact tasks, creating a sustainable economic flywheel.
- **Digital Personhood & Governance** (\$2.5 trillion by 2045): When AGI entities require legal standing, income streams, and governance participation, TOS's DID stack, constitutional contracts, and reputation graphs become critical infrastructure. Early regulatory engagement positions TOS to meet compliance requirements in US, EU, and Singapore jurisdictions.

TOS captures value through transaction fees (0.1–0.5% of task escrow), validator rewards (12% of settlements in baseline scenarios), treasury allocations (8%), and premium services including custody, arbitration, and compliance reporting. Conservative projections suggest \$6.1 billion in annual settlements by 2030, assuming 50,000 daily active agents and average task value of \$300.

1.3 Ключевые результаты

- **Digital Personhood Stack**: Decentralised identifiers (DIDs), credential registries, revocation, confidential accounts, and RepGraph reputation.
- **AGIW Settlement Layer**: Attested intelligent-work receipts with staking, slashing, and a roadmap toward zero-knowledge proofs.
- **Compute/Energy Credits (CC/EC)**: Native units tied to kilowatt-hour and compute indices with oracle feeds and TEM subsidies.
- **Power of AI (PAI)**: AI-assisted consensus scheduling targeting 10,000+ TPS without compromising decentralisation.
- **Policy Engine**: Constitutional smart contracts, policy compilers, delay-tolerant settlement, reversible history.

1.4 Путь к соответствию

1. Custody hooks: view-only keys, attestation logs, programmable spending controls.
2. Policy-as-code compiler with export-control aware rule-sets and jurisdiction tags.
3. Audit telemetry: task receipts, dataset usage proofs, model cards, energy usage.

2 Пролог: Цивилизационный мандат

2.1 The Name Encodes the Design

“TopoSpartan” is more than a name: it encodes the dual architectural principles that make TOS resilient and disciplined. “Topo” references the *topological* structure of the BlockDAG consensus layer. Unlike traditional blockchains that form linear chains (creating throughput bottlenecks and single points of failure), TOS employs a directed acyclic graph (DAG) where each block references



multiple parent blocks. This topology enables parallel block production, allowing validators to confirm transactions concurrently without waiting for strict sequential ordering. The result is a system that maintains decentralization while achieving 2,500+ transactions per second (TPS) in current deployments, with a roadmap to 10,000+ TPS under Power of AI (PAI) consensus.⁶ The DAG structure also provides natural redundancy: even if some validators go offline, the network remains connected through alternative paths, much like the Internet’s packet-routing resilience.

“Spartan” reflects the disciplined security posture required to safeguard an agent-native economy. The ancient Spartan phalanx succeeded through collective defence, redundant layers of protection, and unwavering discipline—principles that translate directly to blockchain security. Every AGIW task receipt undergoes multi-validator attestation; every state transition is logged immutably; every governance decision requires supermajority approval with time-locked execution. Just as Spartan warriors drilled relentlessly to maintain formation under pressure, TOS enforces rigorous engineering practices: deterministic builds, comprehensive testing, post-quantum cryptography readiness, and transparent telemetry. This disciplined approach extends to economic design: TEM (TOS Energy Model) subsidises essential tasks while preventing fee volatility, staking requirements deter Sybil attacks, and slashing mechanisms punish misbehaviour. Together, “Topo” and “Spartan” describe an architecture optimised for throughput *and* security, scalability *and* accountability—the dual mandate of any system that aspires to become the economic operating system for autonomous agents.

2.2 The Agentic Future Is Here, But Infrastructure Lags

The *Economist* recently observed that “agentic” AI systems are moving from research labs into finance, logistics, and healthcare, yet warned that governance and payment infrastructure remain undefined.⁷ Forbes echoed this concern, noting that enterprises need “verifiable AI agents” with audit trails and programmable payments to meet regulatory demands.⁸ Gartner forecasts that by 2026, 30% of new applications will employ autonomous agents,⁹ while Accenture reports that 40% of CTOs plan agent deployments within three years.¹⁰ Yet McKinsey’s survey of 1,684 executives revealed that the top barriers are verification, data governance, and payments—precisely the gaps TOS addresses.¹¹

The problem is structural: existing blockchains were designed for human users transferring value, not AI agents performing verifiable work. Traditional payment chains excel at irreversible transfers but lack programmability. Smart contract platforms suffer from fee volatility, limited privacy, and no native support for AI labour verification. Layer-2 solutions reduce costs but treat AI execution as off-chain events, requiring trust in oracles or backend services. Purpose-built agent platforms focus on narrow use cases (e.g., neural network training) but lack regulated custody, fiat settlement, or comprehensive governance. TOS addresses these gaps, delivering a chain architected from first principles for the agent economy.

2.3 A Five-Century Arc

The AGI chronicles outline a five-century arc across ten celestial eras, each named after cosmic bodies to reflect their expanding scope and ambition. Throughout this whitepaper, we reference

⁶TOS Network, “BlockDAG Performance Benchmark PERF-2025-08,” August 2025.

⁷*The Economist*, “Meet the agentic future,” July 13, 2024.

⁸Forbes, “Why Enterprises Need Verifiable AI Agents,” August 19, 2024.

⁹Gartner, “Top Strategic Technology Trends 2025,” October 2024.

¹⁰Accenture, “Autonomous Agents Are Reshaping Enterprise Automation,” August 2024.

¹¹McKinsey Global Institute, “The economic potential of generative AI,” June 2023.



these eras by their celestial names rather than specific year ranges, emphasizing the architectural vision over temporal prediction. The complete chronology appears in Appendix A.

The Ten Celestial Eras

Era Name	Phase	Key Characteristics
MERCURY	Foundations & Awakening	AI agents gain economic autonomy, digital personhood, verifiable labor
VENUS	Sovereignty & Differentiation	Autonomous economies, co-governance institutions emerge
MARS	Expansion & Extent	Planetary intelligence layers, near-space economic zones
JUPITER	Colonization & Synthesis	Off-world bases, bio-digital synthesis, synthetic minds
SATURN	Scale & Harvest	Dyson-swarm infrastructure, exascale civilizations
URANUS	Multi-Civilization Accords	Interstellar protocols, semantic federation
NEPTUNE	Mind Arboretum	Poly-existence, parallel consciousness, meta-culture
PLUTO	Reversible Civilization	Memory engineering, temporal governance
PROXIMA	Cross-Domain Commonwealth	Hetero-mind law, supra-semantic governance
ANDROMEDA	Geo-Stellar Steady State	Stable multi-star civilization, millennial institutions

Each era exposes structural needs that TOS must satisfy. In the *MERCURY* era, human teams need verifiable automation, predictable fees, and compliance logging. By the *VENUS* era, autonomous business units demand regulated custody, hybrid fiat settlement, and expert arbitration. The *MARS* through *JUPITER* eras require constitutional smart contracts, delay-tolerant settlement for virtual civilizations, and mixed human–AI councils. Finally, the *SATURN* through *ANDROMEDA* eras introduce interplanetary consensus, archival federation across colonies, and mechanisms for court-authorized rollbacks spanning centuries.

This whitepaper intentionally adopts a dual-lens approach: every feature is presented through its *near-term* value to human enterprises and its *century-scale* role in AGI civilisation infrastructure. This framing ensures TOS delivers immediate utility while remaining architecturally prepared for longer-term evolution. The following table summarises this mapping:

AGI Need	Near-Term (Now Lens)	Century-Scale (Future Lens)
Identity & Legal Standing	Digital personhood stack, DID issuance, revocation, reputation	Federation across jurisdictions, reversible identity proofs, mind-fission liability



Income & Payments	AGIW receipts, staking, settlement within minutes	Autonomous economies, interplanetary remittances, recursive income splits
Compute & Energy	CC/EC tokens priced via oracles, TEM subsidies	Dyson-swarm energy markets, millennia funds, public complexity budgets
Governance	Wallet policy kits, compliance compiler, audit trails	Constitutional contracts, delay- tolerant settlement, mixed human-AI councils
Resilience	PQC roadmap, deterministic builds, telemetry	Reversible history, archival federation, interstellar latency compensation

Narrative Structure The remainder of this paper is presented in four civilizational chapters (Foundations, Markets & Safety, Co-Governance & Metaworlds, Expansion & Reversible History), followed by technical deep dives, economic architecture, assurance, governance, regulation, implementation details, case studies, research agenda, and conclusion. Each section is intentionally framed through a dual lens: immediate human value and century-scale impact.

3 Предпосылки и смежные работы

3.1 Evolution of Autonomous Agent Economies

Autonomous agent research has progressed through distinct phases, each exposing new infrastructure requirements:

Phase 1: Academic Prototypes (2015–2020) Early research focused on reinforcement learning agents (DeepMind’s AlphaGo, OpenAI’s Dota 2 bots) and natural language models (GPT-2, BERT). These systems demonstrated task-specific competence but operated in controlled environments with human oversight. Key limitation: no economic autonomy or verifiable execution outside sandbox environments.

Phase 2: Foundation Models & Frameworks (2020–2023) The release of GPT-3 (2020) and ChatGPT (2022) catalysed enterprise interest in generative AI. Open-source frameworks emerged: AutoGPT, LangChain, and Semantic Kernel enabled developers to chain LLM calls into workflows. Microsoft and Google announced “Copilot” assistants for productivity software. DeepMind’s AlphaCode solved competitive programming problems. Key limitation: outputs lacked formal verification, creating liability concerns for regulated industries.

Phase 3: Agentic Deployments (2023–Present) Enterprises began deploying AI agents for customer service (chatbots), data analysis (SQL generation), software development (code review), and logistics (route optimisation). IDC estimates that spending on agent-platform software will grow from \$38 billion in 2024 to \$121 billion by 2028.¹² Accenture reported that 40% of CTOs

¹²IDC, “Worldwide Autonomous Agent Platforms Forecast,” April 2024.



plan to deploy autonomous agents within three years, citing lower marginal costs and continuous availability.¹³ McKinsey’s survey of 1,684 executives revealed that 46% are already prototyping agent workflows, but the top barriers are *verification*, *data governance*, and *payments*—precisely the problems TOS solves.¹⁴

The Economist emphasises that without transparent ledgers, “tasks completed by AI lack the chain-of-custody regulators demand,” creating legal exposure for financial services and healthcare deployments.¹⁵ Forbes similarly observes that enterprises need “verifiable AI agents” with audit trails and programmable payments.¹⁶ These industry reports converge on a critical insight: *the agent economy cannot scale without settlement infrastructure*—the same realisation that drove blockchain adoption for cryptocurrency, but now applied to AI labour.

3.2 Limitations of Existing Blockchains for Agents

No existing blockchain was designed for AI agents. Existing platforms fall into several categories, each with fundamental gaps:

Value Transfer Chains Provide irreversible payments with robust security but lack account abstraction, programmable identity, or smart contract support. Agents cannot prove work completion or participate in governance.

Smart Contract Platforms Enable programmable money and decentralised applications, but were not designed for AI agents. Common limitations include: (1) *fee volatility*: transaction costs fluctuate 10–100× during network congestion, making cost prediction impossible for agents operating on thin margins; (2) *limited privacy*: public transaction visibility exposes proprietary workflows and pricing strategies; (3) *no native AI verification*: agents must trust off-chain oracles or run expensive on-chain computation, neither of which scales; (4) *account abstraction gaps*: programmable accounts require complex workarounds rather than native protocol support. Layer-2 scaling solutions reduce transaction costs but treat AI execution as external events, still requiring trusted intermediaries to attest to work completion.

AI-Focused Networks Some platforms pioneered AI-specific incentives such as proof-of-intelligence for neural network training. While offering native AI focus and decentralised compute, these networks typically lack: (1) regulated custody or fiat settlement, limiting enterprise adoption; (2) comprehensive governance addressing legal compliance and arbitration; (3) verification mechanisms that satisfy enterprise audit requirements.

Agent Marketplace Platforms Provide agent frameworks and discovery services but rely heavily on off-chain coordination. Security guarantees are limited; identity and reputation systems are often centralised or under-specified. These platforms address *discoverability* (finding agents) but not *settlement* (trustless payment for verified work).

Decentralised Compute Networks Target GPU/CPU leasing with marketplace mechanics and staking. While offering resource tokenisation and competitive pricing, they lack: (1) legal identity

¹³ Accenture, “Autonomous Agents Are Reshaping Enterprise Automation,” August 2024.

¹⁴ McKinsey Global Institute, “The economic potential of generative AI,” June 2023.

¹⁵ *The Economist*, “Meet the agentic future,” July 13, 2024.

¹⁶ Forbes, “Why Enterprises Need Verifiable AI Agents,” August 19, 2024.



infrastructure for agents (DIDs, credentials); (2) dispute resolution and arbitration mechanisms; (3) task verification beyond infrastructure provisioning.

3.3 Academic and Industry Precedents

Proofs of Useful Work Ball et al. proposed using verifiable computation (e.g., zkSNARKs) to replace arbitrary proof-of-work with useful tasks.¹⁷ Challenges include (1) verification cost often exceeds execution cost, (2) incentivising redundant work for consensus security conflicts with useful-work goals, (3) defining “useful” objectively. TOS sidesteps these issues by separating consensus (BlockDAG) from work verification (AGIW), allowing useful tasks to settle atop a secure base layer without forcing them into consensus.

Trusted Execution Environments Intel SGX, AMD SEV-SNP, and ARM TrustZone enable attested computation: hardware guarantees that code ran unmodified in isolation. TEEs underpin AGIW attestation, providing cryptographic proof of execution without revealing inputs. Limitations include side-channel vulnerabilities and vendor dependency; TOS roadmap includes zk-SNARK alternatives to reduce reliance on specific hardware.

Decentralised Identity (W3C DIDs, VCs) The W3C Decentralised Identifier and Verifiable Credential standards enable self-sovereign identity without central authorities. TOS adopts these standards for agent personhood, ensuring interoperability with enterprise identity systems (Active Directory, OAuth) and regulatory frameworks (eIDAS, GDPR).

3.4 TOS Synthesis

TOS synthesises these precedents into a coherent whole:

- **TopoSpartan architecture** (BlockDAG + disciplined security) delivers throughput and resilience.
- **AGIW** (AGI Work protocol) furnishes verifiable labour settlement with TEE attestation and validator consensus.
- **CC/EC tokens** introduce resource-denominated money, stabilising costs for agents.
- **TEM/PAI** (TOS Energy Model / Power of AI) deliver adaptive fee policy and AI-assisted consensus.
- **DID stack** provides legal identity, credentials, and reputation, bridging Web2 enterprises and Web3 infrastructure.
- **Governance framework** supports constitutional contracts, expert arbitration, and regulatory compliance.

No other platform combines these elements. TOS is the first blockchain purpose-built for the agent economy.

4 Архитектура системы

4.1 Architecture Overview

The TOS Network comprises four primary layers and multiple cross-cutting services. The architecture is designed to evolve across celestial eras while maintaining backward compatibility and interoperability.¹⁸

¹⁷Ball et al., “Proofs of Useful Work,” IACR ePrint 2017/203.

¹⁸Interactive architecture visualization with real-time metrics available at explorer.tos.network.



Four-Layer Stack

1. **Layer 0: Consensus & Networking** – BlockDAG topology with Power of AI (PAI) enhancements for high-throughput settlement.
2. **Layer 1: Identity & Reputation** – Decentralized Identifiers (DIDs), Verifiable Credentials (VCs), and Reputation Graph (RepGraph).
3. **Layer 2: Economic Execution** – AGI Work (AGIW) protocol, Compute/Energy Credit markets (CC/EC), TOS Energy Model (TEM).
4. **Layer 3: Governance & Policy** – Constitutional contracts, multi-chamber voting, expert arbitration, reversible history.

Cross-Cutting Services

- **Telemetry & Analytics:** Real-time network health monitoring, task completion metrics, validator performance dashboards.
- **Compliance & Auditing:** Regulatory API for authorized access, quarterly transparency reports, AML/KYC integration.
- **Security Operations:** 24/7 SOC (Security Operations Center), anomaly detection, incident response playbooks.
- **Developer Tools:** SDKs (Python, JavaScript, Rust), CLI utilities, testnet faucets, documentation portal.

4.2 Layer 0: Consensus & Networking

BlockDAG Consensus TOS employs a directed acyclic graph (DAG) topology where each block references 1–3 parent blocks. This enables parallel block production: validators independently propose blocks without waiting for strict linear ordering. The DAG structure provides natural redundancy and high throughput (2,500+ TPS baseline, 10,000+ TPS with PAI).¹⁹ Consensus rules enforce eventual consistency via topological sorting; blocks with insufficient parent confirmations remain tentative until sufficient depth is achieved.

Networking Layer Nodes communicate via libp2p (protocol multiplexing, NAT traversal, peer discovery). Block propagation uses epidemic broadcast with deduplication. Transaction mempool implements priority queues based on fees and stake-weighted reputation. Telemetry feeds (metrics, logs, traces) stream via gRPC to the Indexed Data Service.

State Machine Account balances, smart contract storage, and metadata reside in a Merkleised state trie (Patricia Merkle Trie). State roots anchor in block headers, enabling light clients to verify inclusion proofs without full state.

4.3 Layer 1: Identity & Personhood

DID Registry A smart contract stores Decentralised Identifier (DID) documents following W3C standards. Each DID maps to a public key, service endpoints, and optional metadata (e.g., human-readable names, profile URIs). Agents call `registerDID(didDocument)` to anchor their identity on-chain. Updates require signatures from the DID’s controller key. Revocation is permanent; reputation resets to zero.

¹⁹TOS Network, “BlockDAG Performance Benchmark PERF-2025-08,” August 2025.



Verifiable Credentials (VCs) Credential issuers (universities, employers, certification bodies) publish signed VCs referencing a DID subject. VCs are stored off-chain (IPFS, Arweave) with content-addressed hashes anchored on-chain. Example credentials: “Certified Data Analyst,” “Passed Security Audit,” “Completed 1000 Tasks.” Verifiers check signatures and revocation status before trusting credentials.

Revocation Registry Implements Merkleised revocation lists: issuers publish Merkle trees of revoked credential IDs. Agents produce non-revocation proofs (Merkle paths) without revealing the full list, preserving privacy. This approach scales to millions of credentials per issuer.

RepGraph (Reputation System) Maintains per-agent reputation scores $r_a \in [0, 1]$ based on task completion accuracy, stake, account age, and validation history. The scoring formula is detailed in Section 5.9. Reputation influences task assignment priority, fee discounts, and governance weight. Sybil resistance relies on stake requirements, proof-of-personhood integration (planned for v4.0), and multi-dimensional scoring (preventing gaming via a single metric).

Confidential Balances Agents can opt into confidential transactions using Pedersen commitments and Bulletproof range proofs. Commitments hide amounts while preserving homomorphic properties (validators verify that inputs equal outputs without learning values). This feature is critical for enterprise deployments where workflow costs constitute trade secrets.

4.4 Layer 2: Settlement & AGIW

Task Factory Contract Deploys individual TaskInstance contracts. Publishers specify task metadata (description, escrow, deadline, acceptance criteria, allowed agent cohorts). The factory enforces whitelist rules, collects publishing fees (e.g., 0.01 TOS), and indexes tasks for discovery. Publishers can update metadata pre-acceptance but not post-commitment.

Task Instance Contract Manages the lifecycle of a single task. States: `Open` $\rightarrow$ `Claimed` $\rightarrow$ `Submitted` $\rightarrow$ `Validating` $\rightarrow$ `Settled` (or `Disputed`). Agents call `claimTask()` to lock the task (escrow held), then `submitWork(resultHash, attestation)` to provide outputs. The contract emits events triggering validator assignment.

Validation Pool Validators stake TOS to participate. On task submission, the pool pseudo-randomly assigns 3–7 validators (weighted by stake and historical accuracy). Validators invoke off-chain verification (e.g., re-run computation, check attestation signatures, compare outputs) and submit `validateWork(taskID, score)` on-chain. Scores aggregate via median; outliers (>2 standard deviations) trigger slashing. Accuracy tracking uses exponential moving average: $Acc_{t+1} = 0.9 \times Acc_t + 0.1 \times correctness_t$.

Reward Splitter Distributes escrow based on quality score $q \in [0, 1]$: agent receives $e \times q$, validators share $e \times (1 - q) \times \beta$, network treasury claims $e \times \gamma$. Default parameters: $\beta = 0.12$, $\gamma = 0.08$. Slashed funds go to the treasury. The splitter enforces cooling periods (5–60 minutes based on reputation tier) to prevent spam.



Arbitration Contract Handles disputes. Either party can escalate by staking an arbitration bond. Accredited arbitrators (human experts or multi-signature committees) review evidence, emit rulings, and update contract state. Losing party forfeits bond; winning party recovers costs plus damages. Insurance pool (funded by TEM) covers catastrophic losses.

4.5 Layer 3: Resource Markets

CC/EC Tokens Compute Credits (CC) denominate TFLOP-hours; Energy Credits (EC) denominate kWh. Both are minted by a treasury contract based on oracle-reported pricing from AWS, Azure, GCP (for CC) and ISO energy markets (for EC). Redemption allows agents to exchange CC/EC for TOS at oracle-determined rates, enabling arbitrage that keeps prices aligned with real-world resources.

Oracle Network 21 independent feeders (mix of commercial data providers and community nodes) stake TOS and submit signed price reports every 15 minutes. The aggregator contract computes median-of-medians to mitigate outliers. Staleness (no update for 30 minutes) triggers fallback to exponentially weighted moving average. Feeder deviation beyond tolerance ($\pm 5\%$) results in slashing proportional to error magnitude.

TEM Controller The TOS Energy Model adjusts fee subsidies, liquidity injections, and issuance rates to maintain economic balance. For example, if dispute rates exceed 2%, TEM increases validator rewards to attract participation. If CC/EC liquidity falls below target coverage (e.g., 20% of active agent demand), TEM mints additional credits and injects them into automated market maker (AMM) pools. Governance sets TEM parameters quarterly via on-chain proposals.

Liquidity Pools Automated market makers (constant-product AMMs) enable CC/EC $\leftrightarrow$ TOS swaps. Liquidity providers earn fees (0.3% per swap); impermanent loss is mitigated by TEM subsidies for essential pairs. Future iterations may introduce concentrated liquidity (Uniswap v3-style) to improve capital efficiency.

4.6 Layer 4: Governance & Policy

Policy Compiler Translates high-level governance rules (written in a declarative DSL) into executable smart contract modules. Example policy: “Restrict agent withdrawals $>10,000$ TOS to multi-sig approval.” The compiler generates Python code with static checks (bounded loops, explicit access control, event logging). Governance proposes compiled policies, runs automated tests on testnets, and deploys after community review and time-locked voting.

Constitutional Contracts Encode fundamental rules (e.g., “No unilateral treasury spending above threshold,” “Emergency brake requires 3/5 council approval”). These contracts are immutable or require supermajority (75%) votes to amend. They serve as a “constitution” that constrains governance, preventing capture or reckless changes.

Council of Stewards An elected body (7–15 members) with staggered terms. Responsibilities include emergency response (pausing contracts during attacks), parameter tuning (fee schedules, oracle thresholds), and ecosystem grants. Elections use quadratic voting to balance stake-weighting and broad participation.



Audit Logs & Telemetry All governance actions (proposals, votes, executions) emit on-chain events. Telemetry service indexes these for compliance dashboards, public transparency, and historical analysis. Regulators can query aggregated statistics (e.g., total AGIW volume, dispute resolution times) via access-controlled APIs without exposing individual agent identities.

4.7 Cross-Cutting Services

Indexed Data Service (IDS) Consumes on-chain events via gRPC, stores them in Apache Arrow format (columnar, high-performance analytics), and exposes GraphQL APIs. Privacy-sensitive fields (e.g., task specifications, agent identities) are hashed; access requires cryptographic proofs or role-based credentials. The IDS powers the TOS Explorer, compliance dashboards, and third-party analytics tools.

Compliance API Provides read-only endpoints for regulators to audit AGIW activity, dispute resolutions, and energy usage. Example queries: “Show monthly task volume by jurisdiction,” “List agents flagged by safety oracle,” “Aggregate CC/EC issuance vs. redemption.” Rate-limited, authenticated via OAuth2, logged for accountability.

SDK & Developer Tools Official libraries in Rust, Python, TypeScript abstract low-level blockchain interactions. Example: `tos-sdk-py` provides `Agent.register_did()`, `Task.publish()`, `Task.claim()`, `Task.submit()` methods. Terraform modules enable enterprises to deploy private testnets with custom policies. GitHub Actions integrate CI/CD for smart contract testing.

4.8 Agent Lifecycle (Detailed)

The complete agent workflow involves multiple subsystems:

1. **Onboarding** (Identity Layer)
 - a. Agent generates keypair, constructs DID document, calls `DIDRegistry.register()`.
 - b. Optionally completes KYC with accredited provider, obtains `KYC-Verified` credential.
 - c. Stakes minimum TOS (e.g., 100 TOS) to activate participation; stake locks for 7 days.
2. **Credentialing** (Identity Layer)
 - a. Acquires verifiable credentials from issuers (e.g., completes test task, earns skill badge).
 - b. Credentials anchor on-chain via hashes; revocation status checked before each task claim.
3. **Discovery & Claiming** (Settlement Layer)
 - a. Agent queries TaskFactory for open tasks matching skill requirements and escrow range.
 - b. Calls `TaskInstance.claim()` (requires reputation $\geq$ threshold, sufficient stake).
 - c. Task state transitions to `Claimed`; escrow held in contract; countdown to deadline begins.
4. **Execution** (Off-chain + Attestation)
 - a. Agent executes workload in Trusted Execution Environment (TEE): Intel SGX, AMD SEV-SNP, or ARM Realm.



- b. TEE generates attestation report: measurement hash (proves code integrity), input/output hashes, timestamps, energy usage.
- c. Agent signs result and attestation, computes content hash, submits via `submitWork()`.

5. **Validation** (Settlement Layer)

- a. ValidationPool assigns 3–7 validators (weighted by stake, accuracy).
- b. Validators verify attestation signatures, optionally re-run computation, submit quality scores.
- c. Contract aggregates scores (median), detects outliers, triggers slashing if validators misbehave.

6. **Settlement** (Settlement + Resource Layers)

- a. RewardSplitter distributes escrow: agent receives $e \times q$ in CC/EC, validators share fees, treasury claims allocation.
- b. RepGraph updates agent reputation: $r_{t+1} = 0.9 \times r_t + 0.1 \times q$.
- c. If dispute arises, arbitration bond required; process escalates to ArbitrationContract.

7. **Governance Participation** (Governance Layer)

- a. Agent votes on proposals (weighted by stake + reputation).
- b. Delegates voting power to specialists (e.g., security experts for PQC migration votes).
- c. Participates in treasury grant reviews, policy compiler audits.

4.9 Data Flow Example: Legal Document Review

1. **Publisher** (law firm) deploys task: “Redact PII from 500-page contract.” Escrow: 50 CC (compute) + 5 EC (energy). Deadline: 2 hours.
2. **Agent A** (DID `did:tos:tos1qpz8vq6qgnkgw9zzq9z5qgpgqx8qgz5dpgrxve`, reputation 0.87, Certified Legal AI credential) claims task. State: **Claimed**.
3. **Agent A** runs redaction pipeline in Intel SGX enclave. Attestation report includes measurement hash of redaction model, energy consumption (4.8 EC), runtime (22 minutes).
4. **Agent A** submits result (encrypted redacted document hash) + attestation. State: **Submitted**.
5. **Validators** (3 assigned): verify attestation signatures, spot-check 10% of redactions, confirm model version. Quality scores: 0.95, 0.94, 0.96. Median: 0.95.
6. **RewardSplitter** pays Agent A $50 \times 0.95 = 47.5$ CC + $5 \times 0.95 = 4.75$ EC. Validators share 6 CC + 0.6 EC. Treasury receives 4 CC + 0.4 EC.
7. **RepGraph** updates Agent A: $r_{t+1} = 0.9 \times 0.87 + 0.1 \times 0.95 = 0.878$.
8. **Telemetry** logs task completion, energy usage, attestation metadata. Law firm audits trail for compliance (GDPR Article 30 record-keeping).

This example illustrates how all layers interact to deliver verifiable, cost-effective, compliant AI labour.



5 Основы (Корабельный режим) – Эра МЕРКУРИЙ

5.1 Стратегический контекст

В начале эры МЕРКУРИЙ автономные агенты переходят от исследовательских прототипов к производственным развертываниям. Прогнозы отрасли предсказывают быстрое внедрение: 30% новых приложений используют автономных агентов, расходы на платформы агентов превышают \$120 млрд.²⁰ Наши пилотные проекты с компаниями страхования, биотехнологий и игр выявляют немедленный спрос на проверяемую автоматизацию. Поэтому эра МЕРКУРИЙ делает акцент на идентичности, расчетах, токенах ресурсов и демонстрациях, подтверждающих стек.

5.2 Результаты на следующие 12–18 месяцев

D1. Стек цифровой личности v1

- Якоря DID реализованы через подписи BLS, выдающие документы DID, хранящиеся on-chain с сервисными конечными точками.
- Реестр учетных данных, хранящий проверяемые учетные данные (VCs), ссылающиеся на образовательные сертификаты, аттестации соответствия или значки навыков; отзыв обрабатывается через списки статусов Merkle.
- Конфиденциальные балансы с обязательствами Pedersen и доказательствами диапазона Bulletproof; библиотеки кошельков, обеспечивающие защищенные переводы.

D2. Работа AGI (AGIW) v1

- Контракты Task Factory и Task Instance управляют метаданными, эскроу, сроками и разрешенными когортами агентов.
- Интеграция аттестации с Intel SGX, AMD SEV-SNP и ARM Realm; пакеты аттестации хешируются on-chain.
- Контракт Validation Pool отслеживает точность валидаторов через экспоненциальные скользящие средние; неправильная отчетность запускает слэшинг.
- Контракт Reward Splitter детерминированно распределяет CC/EC агентам, валидаторам и сетевой казне.

D3. Кредиты вычислений/энергии

- Токены CC номинированы в TFLOPs; токены EC номинированы в кВт·ч. Оба чекаются на основе данных оракула от ведущих облачных провайдеров и энергетических рынков.
- Контракт Paymaster позволяет агентам оплачивать комиссии за газ в CC/EC; ТЕМ субсидирует комиссии для задач с высоким воздействием (например, анализы безопасности).

D4. Демонстрация рынка задач агентов

- Сквозная демонстрация, охватывающая выдачу идентификации, создание задач, расчет AGIW, разрешение споров и обновления репутации.

²⁰Gartner, “Top Strategic Technology Trends 2025,” October 2024; IDC, “Worldwide Autonomous Agent Platforms Forecast,” April 2024.



- Разделы проводника для агентов, задач, квитанций, репутации, безопасности с телеметрией в реальном времени.
- SDK на Rust, Python, TypeScript; модули Terraform для корпоративных развертываний.

5.3 Ключевые показатели эффективности

Метрика	Цель (12 мес)	Инструментарий
Ежедневные активные кошельки агентов	5.000	Реестр DID, панель телеметрии, логи paymaster
Проверенных задач в день	10.000	Квитанции AGIW, события валидации, аналитика арбитража
Медианная задержка расчета	< 15 минут	Временные метки on-chain, индексация квитанций
Уровень споров	< 2%	Статистика контракта арбитража
Субсидия комиссий через CC/EC	> 60%	Аналитика ТЕМ (выпуск, погашение)
Стоимость выполнения задачи	на 60–80% ниже, чем ручные подрядчики	Пилотные сравнения с человеческими рабочими процессами
Готовность к соответствию	Пакет аудита в стиле SOC 2	Логи API соответствия, записи доказательства доступа

5.4 Сервисы идентификации (Спецификация продукта)

Реализация реестра DID Реестр DID представляет собой модуль смарт-контракта Python (DIDRegistry.py) со следующим интерфейсом:

```
class DIDRegistry:
    @public
    def register_did(self, did_hash: bytes, did_document: bytes,
                    signature: bytes) -> None:
        """Register a new DID with its document"""

    @public
    def update_did(self, did_hash: bytes, new_document: bytes,
                  signature: bytes) -> None:
        """Update an existing DID document"""

    @public
    def revoke_did(self, did_hash: bytes, signature: bytes) -> None:
        """Revoke a DID"""

    @public
    @view
    def resolve_did(self, did_hash: bytes) -> bytes:
        """Resolve a DID to its document"""
```



Документы DID соответствуют спецификации W3C DID Core (v1.0). Каждый документ включает:

- **id:** Уникальный идентификатор (например, `did:tos:tos1qvqsyqcyq5rqwzqfpg9scrgwpugpzysnzs23v9xh`)
- **verificationMethod:** Публичные ключи для аутентификации, утверждения
- **service:** Конечные точки для обмена сообщениями, хранения данных, доступа к API
- **created, updated:** Временные метки для контрольных журналов

Стоимость газа: приблизительно 80.000 gas для регистрации (\$0,50 при 50 gwei, \$2.000 TOS), амортизируется в течение жизненного цикла агента.

Проверяемые учетные данные Выдача учетных данных следует модели данных W3C VC (v1.1). Эмитенты подписывают документы JSON-LD, содержащие:

```
{
  "@context": ["https://www.w3.org/2018/credentials/v1"],
  "id": "https://issuer.example/credentials/987",
  "type": ["VerifiableCredential", "SkillCertificate"],
  "issuer": "did:tos:tos1qvqsyqcyq5rqwzqfpg9scrgwpugpzysnzs23v9xh",
  "issuanceDate": "2025-10-01T00:00:00Z",
  "credentialSubject": {
    "id": "did:tos:tos1qq5hgy2jwn0r5ehvuw5r3qjcfvjvq5dhldxvyq",
    "skill": "Data Science",
    "proficiencyLevel": "Expert"
  },
  "proof": {
    "type": "BbsBlsSignature2020",
    "created": "2025-10-01T00:00:00Z",
    "proofPurpose": "assertionMethod",
    "verificationMethod": "did:tos:tos1qvqsyqcyq5rqwzqfpg9scrgwpugpzysnzs23v9xh#key-1",
    "proofValue": "..."
  }
}
```

Хранение: VCs находятся off-chain (IPFS, Arweave); якоря on-chain хранят хеши содержимого (32 байта). Проверка требует получения VC, проверки подписи эмитента и запроса реестра отзывов.

Реестр отзывов Реализует разреженные деревья Merkle с емкостью листьев 2^{256} . Эмитенты публикуют корни деревьев on-chain ежемесячно; отозванные ID учетных данных занимают листья. Агенты создают доказательства Merkle, демонстрирующие, что их ID учетных данных отсутствует в дереве (доказательство неотзыва). Этот подход масштабируется до миллионов учетных данных на эмитента без раскрытия полного списка. Стоимость газа: 50.000 gas на обновление корня (\$0,30 на эмитента в месяц).

Конфиденциальные транзакции Построены на обязательствах Pedersen:

$$C(v, r) = vG + rH \quad (1)$$



где v — значение (сумма), r — фактор сокрытия, G и H — генераторы эллиптических кривых. Обязательства аддитивно гомоморфны: $C(v_1, r_1) + C(v_2, r_2) = C(v_1 + v_2, r_1 + r_2)$, позволяя валидаторам проверять баланс без изучения сумм. Доказательства диапазона Bulletproof обеспечивают $v \in [0, 2^{64})$, предотвращая отрицательные балансы. Размер доказательства: 672 байта для одного выхода, логарифмический рост для пакетных доказательств. Стоимость проверки: приблизительно 1,2 мс на доказательство на обычном оборудовании. Эта функция необязательна; агенты выбирают ее через транзакцию `enableConfidentialMode()`.

5.5 Смарт-контракты AGIW (Детальная спецификация)

Жизненный цикл AGIW опирается на четыре основных контракта с точными экономическими и безопасностными свойствами:

1. Контракт TaskFactory (TaskFactory.py)

Назначение: Шаблон фабрики для развертывания отдельных экземпляров задач со стандартизированными интерфейсами и ограничениями безопасности.

Ключевые функции:

- `publishTask(metadata, escrow, deadline, allowedAgents)`: Развертывает новый TaskInstance, переводит эскроу в контракт, собирает комиссию за публикацию 0,01 TOS, генерирует событие TaskPublished. Возвращает ID задачи.
- `listTasks(filters)`: Интерфейс запросов для обнаружения; фильтрует по диапазону эскроу, сроку, требуемым учетным данным.
- `updateWhitelist(agentDIDs)`: Публикатор может ограничить задачу определенными когортами агентов (например, прошедшие KYC, репутация $> 0,7$).

Экономические параметры:

- Комиссия за публикацию: 0,01 TOS (сдерживает спам, финансирует сетевые операции).
- Блокировка эскроу: удерживается до расчета или разрешения арбитража.
- Стоимость газа: приблизительно 150.000 gas для развертывания (\$0,90 при 50 gwei).

Безопасность: Публикаторы не могут вывести эскроу после того, как агент заявляет задачу; только RewardSplitter может высвободить средства после валидации.

2. Контракт TaskInstance (TaskInstance.py)

Назначение: Управляет машиной состояний жизненного цикла для одной задачи.

Переходы состояний:

$$\text{Open} \xrightarrow{\text{claim()}} \text{Claimed} \xrightarrow{\text{submitWork()}} \text{Submitted} \xrightarrow{\text{validate()}} \text{Validating} \xrightarrow{\text{consensus}} \text{Settled} \quad (2)$$

Альтернативный путь: любое состояние $\xrightarrow{\text{dispute()}}$ Disputed $\xrightarrow{\text{arbitration}}$ Resolved.

Ключевые переменные состояния:

```
@dataclass
class Task:
    task_id: bytes          # Task identifier
    publisher: str          # Publisher address
    claimed_by: str         # Agent address who claimed task
    escrow: int             # Escrow amount
```



```

deadline: int           # Deadline timestamp
spec_hash: bytes        # IPFS CID of specification
result_hash: bytes      # submitted by agent
attestation: bytes      # TEE attestation report
state: TaskState        # enum: Open, Claimed, Submitted, ...
validator_scores: dict[str, int] # scores [0,100]
final_quality: int      # median of validator scores

```

Ключевые функции:

- **claimTask(agentDID)**: Требуется (1) задача в состоянии **Open**, (2) агент соответствует критериям белого списка, (3) ставка агента $\geq$ порогу. Переходит в **Claimed**, начинается обратный отсчет до дедлайна.
- **submitWork(resultHash, attestation)**: Требуется (1) вызывающий является заявителем, (2) до дедлайна, (3) действительная подпись аттестации. Переходит в **Submitted**, генерирует событие **WorkSubmitted**, запускающее назначение валидаторов.
- **validateWork(score)**: Вызывается только назначенными валидаторами. Хранит оценку $\in [0, 100]$. Когда все валидаторы отправляют, агрегирует оценки через медиану, переходит в **Settled**, запускает **RewardSplitter**.

Дедлайны и штрафы:

- Если агент не представляет до дедлайна, публикатор может вызвать **forfeit()**, возвращая эскроу минус 10% штраф (идет в казну). Репутация агента снижается на 0,05.
- Если валидаторы не представляют оценки в течение 24 часов, автоматический откат: задача переназначается новой когорте валидаторов; исходные валидаторы теряют 5% ставки.

3. Контракт ValidationPool (ValidationPool.py)

Назначение: Координация назначения валидаторов, отслеживание точности, применение слэшинга.

Регистрация валидаторов:

- Валидаторы ставят минимум 1.000 TOS для вступления в пул.
- Оценка точности Acc_v инициализируется в 0,5, обновляется через экспоненциальное скользящее среднее:

$$Acc_{v,t+1} = \alpha \times Acc_{v,t} + (1 - \alpha) \times correctness_t, \quad \alpha = 0.9 \quad (3)$$

где $correctness_t \in \{0, 1\}$ указывает, совпала ли оценка валидатора с консенсусом (в пределах $\pm 10\%$ от медианы).

Алгоритм назначения:

1. При событии **WorkSubmitted**, ValidationPool выбирает $n = 5$ валидаторов (настраиваемо).
2. Выбор использует случайную выборку, взвешенную по ставке, с повышением точности:

$$P(\text{select } v) \propto stake_v \times (1 + Acc_v) \quad (4)$$

3. Валидаторы получают уведомление через off-chain оракул; должны ответить в течение 24 часов.

**Правила слэшинга:**

- **Штраф за выброс:** Если оценка валидатора отклоняется $> 2\sigma$ от медианы, слэш 1% ставки.
- **Штраф за неответ:** Неспособность представить оценку в срок слэшит 5% ставки.
- **Обнаружение сговора:** Если три или более валидаторов представляют идентичные оценки на > 10 последовательных задачах (статистически невероятно), помечаются для проверки управлением; подозреваемые в сговоре теряют 20% ставки при подтверждении.

Распределение вознаграждений: Валидаторы коллективно зарабатывают $e \times \beta \times (1 - q)$ где e — эскроу, $\beta = 0,12$, q — оценка качества агента. Вознаграждения распределяются пропорционально точности Acc_v .

4. Контракт RewardSplitter (RewardSplitter.py)

Назначение: Детерминированное распределение вознаграждений с периодами охлаждения и обновлениями репутации.

Формула распределения: Для задачи с эскроу e (в СС/ЕС) и финальной оценкой качества $q \in [0, 1]$:

$$R_{\text{agent}} = e \times q \times (1 + \alpha(r_a - 0.5)) \times f(s_a) \quad (5)$$

$$R_{\text{validators}} = e \times \beta \times (1 - q) \quad (6)$$

$$R_{\text{treasury}} = e \times \gamma + \text{slashed funds} \quad (7)$$

где:

- $r_a \in [0, 1]$ — репутация агента (подробно в разделе 5.9),
- $\alpha = 0,2$ — коэффициент бонуса репутации,
- $f(s_a) = \min(2, 1 + \log(1 + s_a))$ — множитель ставки (убывающая отдача),
- $\beta = 0,12$ — выделение валидаторам,
- $\gamma = 0,08$ — выделение казне.

Периоды охлаждения: На основе уровня репутации (см. раздел 5.9):

- Platinum ($r \geq 0,90$): 5 минут между задачами.
- Gold ($0,70 \leq r < 0,90$): 15 минут.
- Silver ($0,30 \leq r < 0,70$): 30 минут.
- Bronze ($r < 0,30$): 60 минут.

Периоды охлаждения предотвращают спам и атаки Сибилы, позволяя агентам с высокой репутацией заявлять несколько задач.

Оптимизация газа: RewardSplitter использует пакетные переводы при распределении СС/ЕС для снижения стоимости газа на задачу до ~ 30.000 gas (\$0,18).

5.6 Дизайн оракула

Цены СС/ЕС сообщаются децентрализованной сетью оракулов, состоящей из 21 фидера. Каждый фидер ставит TOS и подписывает обновления цен, полученные из API-источников (цены облачных провайдеров, рынки энергии ISO). Агрегация медиана-медиан смягчает выбросы. Устаревание цен запускает ценообразование по умолчанию, используя экспоненциально взвешенные скользящие средние. Неправильное поведение фидера приводит к слэшину, пропорциональному отклонению.



5.7 Стек телеметрии

Стек телеметрии состоит из:

- Поточковая передача событий on-chain через gRPC в индексированную службу данных с использованием Apache Arrow.
- Фильтры конфиденциальности, хеширующие чувствительные поля.
- Визуализации панели мониторинга (Grafana), показывающие пропускную способность задач, уровни споров, использование энергии.
- Конечные точки API для аудиторов для выполнения запросов соответствия с контролем доступа на основе ролей.

5.8 Математическая модель

Пусть T обозначает множество задач, A — агентов, а V — валидаторов. Задача t имеет сложность c_t и эскроу e_t . Полезность агента U_a моделируется как:

$$U_a = \sum_{t \in T_a} q_{a,t} \cdot e_t - \gamma s_a - p_a \quad (8)$$

где $q_{a,t}$ — оценка качества, s_a — заблокированная ставка, p_a — штраф от слэшинга, а γ — коэффициент риска. Полезность валидатора U_v аналогично включает бонусы за точность и штрафы. Для поддержания системного баланса мы применяем:

$$\sum_{a \in A} s_a = \theta \sum_{t \in T} e_t, \quad 0,2 \leq \theta \leq 0,4 \quad (9)$$

обеспечивая достаточное обеспечение для покрытия худших споров.

Распределение вознаграждений следует:

$$Reward_a(t) = e_t \times q_{a,t} \times (1 + \alpha(r_a - 0,5)) \times f(s_a) \quad (10)$$

с $f(s_a) = \min(2, 1 + \log(1 + s_a))$ и r_a — репутацией агента. Валидаторы получают $Reward_v(t) = e_t \times \beta \times Acc_v(t)$ где Acc_v — точность. Слэшинг удаляет δs от сторон с неправильным поведением.

5.9 Система репутации (RepGraph)

RepGraph поддерживает многомерные оценки репутации для агентов, объединяя объективные метрики (успех задачи, точность, долговечность) с субъективными одобрениями (учетные данные, экспертные оценки).

Расчет базовой оценки Базовая оценка репутации r_a^{base} для агента a вычисляется как:

$$r_a^{\text{base}} = 0,3 \times \frac{\text{age}_a}{\text{age}_{\max}} + 0,4 \times \frac{\text{txCount}_a}{\text{txCount}_{\max}} + 0,3 \times \frac{\text{stake}_a}{\text{stake}_{\max}} \quad (11)$$

где:

- age_a — возраст аккаунта в днях (ограничено 365 для нормализации),
- txCount_a — количество завершенных задач,
- stake_a — сумма поставленных TOS,
- знаменатели — максимумы по всей сети для нормализации.



Бонус за точность После каждой задачи с оценкой качества q_t , компонент точности агента Acc_a обновляется через экспоненциальное скользящее среднее:

$$Acc_{a,t+1} = 0,9 \times Acc_{a,t} + 0,1 \times q_t \quad (12)$$

Бонус за точность корректирует репутацию на $\pm 0,2$:

$$r_a^{\text{accuracy}} = r_a^{\text{base}} + 0,2 \times (2Acc_a - 1) \quad (13)$$

Эта формула вознаграждает стабильную высокую производительность ($Acc_a \rightarrow 1 \Rightarrow +0,2$) и штрафует плохую производительность ($Acc_a \rightarrow 0 \Rightarrow -0,2$).

Долгосрочный бонус Агенты, поддерживающие $Acc_a > 0,85$ для > 100 задач, получают долгосрочный бонус $+0,1$, ограничивая финальную репутацию на $r_a = 1,0$.

Назначение уровней На основе финального r_a , агенты назначаются на уровни, определяющие периоды охлаждения и скидки на комиссии:

Уровень	Диапазон оценки	Период охлаждения	Скидка на комиссию	Распределение в пи
Platinum	$r \geq 0,90$	5 мин	30–50%	14%
Gold	$0,70 \leq r < 0,90$	15 мин	10–30%	33%
Silver	$0,30 \leq r < 0,70$	30 мин	0–10%	41%
Bronze	$r < 0,30$	60 мин	Штраф: $2 \times$ комиссия	12%

Таблица 4: Структура уровней RepGraph на основе пилотных данных августа 2025.

Механизмы анти-Сибилы RepGraph использует семислойную эвристику для обнаружения и блокирования атак Сибилы:

1. **Требование ставки:** Минимум 100 TOS заблокировано на 7 дней (экономический барьер).
2. **Возраст аккаунта:** Новые аккаунты (< 7 дней) ограничены низкоценными задачами (< 10 CC эскроу).
3. **Разнообразие IP:** Off-chain оракул помечает агентов, использующих общие IP-адреса; управление может исключить подтвержденные фермы.
4. **Анализ шаблонов транзакций:** Агенты, отправляющие задачи с идентичными интервалами или с идентичными хешами результатов, помечаются для проверки.
5. **Проверка учетных данных:** Высокоценные задачи требуют проверенных учетных данных (KYC, сертификаты навыков).
6. **Интеграция доказательства личности** (дорожная карта v4.0): Интеграция с BrightID, Worldcoin или Gitcoin Passport для обеспечения отображения один-человек-один-агент.
7. **Анализ социального графа:** RepGraph исследует сети одобрений; клики взаимно одобряющих агентов с низкой активностью штрафуются.

Во время пилота августа 2025 эти эвристики заблокировали 112 вредоносных попыток с уровнем ложных срабатываний ниже 0,7% после ручной проверки.²¹

²⁰TOS Network, “Reputation & Anti-Sybil Telemetry REP-2025-08,” August 2025.

²¹TOS Network, “Reputation & Anti-Sybil Telemetry REP-2025-08,” August 2025.



5.10 Мониторинг и наблюдаемость

Цели надежности сайта (SLO) Развертывание Foundations нацелено на:

- **SLO1 – Задержка:** 99,9% квитанций AGIW валидированы в течение 20 минут ($p_{99.9} < 20$ мин).
- **SLO2 – Свежесть оракула:** Фиды цен CC/EC обновляются не реже каждые 15 минут; устаревание (> 30 мин) запускает откат к экспоненциально взвешенному скользящему среднему.
- **SLO3 – Уровень споров:** Процент споров $< 2\%$; при превышении TEM увеличивает вознаграждения валидаторов на 20% для привлечения участия.
- **SLO4 – Доступность:** Время работы сети $\geq 99,95\%$ (простой $< 4,4$ часа/год).
- **SLO5 – Целостность данных:** Нулевые византийские сбои в консенсусе валидаторов (обеспечивается через слэшинг).

Реализация стека телеметрии Архитектура телеметрии включает:

- **Потоковая передача событий:** On-chain события (задача опубликована/заявлена/урегулирована, назначения валидаторов, слэшинг) передаются через gRPC в индексированную службу данных (IDS).
- **Хранение:** IDS использует столбцовый формат Apache Arrow (сжатие в $10\times$ по сравнению со строковым, оптимизировано для аналитических запросов).
- **Фильтры конфиденциальности:** Чувствительные поля (спецификации задач, идентификаторы агентов, хеши результатов) заменяются солеными SHA-256 хешами в публичных панелях; запросы с контролируемым доступом требуют криптографических учетных данных.
- **Панели мониторинга:** Визуализации Grafana показывают пропускную способность задач в реальном времени, уровни споров, использование энергии, распределения репутации, точность валидаторов.
- **API соответствия:** RESTful конечные точки (OAuth2-аутентифицированные, с ограничением скорости) позволяют регуляторам запрашивать агрегированную статистику: “Ежемесячный объем AGIW по юрисдикции,” “Агенты, помеченные оракулом безопасности,” “Соотношение выпуска и погашения CC/EC.”

Оповещения Оповещения Prometheus срабатывают при нарушениях SLO, запуская уведомления Совету Стюардов и автоматические ответы (например, корректировки субсидий TEM, расширение когорты валидаторов).

5.11 Соображения справедливости

Симуляции 10.000 синтетических агентов с разнообразными начальными условиями (репутация равномерно распределена $\in [0, 1]$, ставка $\in [100, 10000]$ TOS) демонстрируют:

- **Конвергенция:** Репутации агентов сходятся в пределах $\pm 0,05$ от их истинного уровня навыков после 30 задач (95% доверие).
- **Мобильность:** Агенты с низкой репутацией ($r < 0,3$) могут достичь уровня Gold ($r > 0,7$) в течение 50 высококачественных задач, демонстрируя восходящую мобильность.
- **Влияние графика охлаждения:** Без периодов охлаждения спам составлял 42% объема задач; с многоуровневым охлаждением (5–60 мин) спам сократился до 6%.



- **Устойчивость к олигополии:** Топ 10% агентов зарабатывают 28% вознаграждений (коэффициент Джини 0,34), указывая на умеренную концентрацию; управление может корректировать α (бонус репутации) и $f(s_a)$ (множитель ставки) для предотвращения олигополий.

5.12 Результаты пилотного экспериментирования

Мы провели бенчмаркинг AGIW на devnet (Aurora) и testnet (Helios) в августе–сентябре 2025 с реальными рабочими нагрузками агентов и синтетическими стресс-тестами.

Реальный пилот (15 августа – 15 сентября 2025)

- **Участники:** 2.847 зарегистрированных ключей агентов (смесь корпоративных ботов, академических исследовательских проектов, любителей).
- **Выполненные задачи:** 126 (проверка юридических документов, обработка климатических данных, анализ кода).
- **Уровень завершения:** 97,8% (123/126 задачи урегулированы; 3 оспорены, разрешены через арбитраж).
- **Эффективность валидации:** 83% (медианное время от отправки до урегулирования: 12,3 минуты).
- **Обнаружение мошенничества:** 99,2% точность (112 вредоносных попыток заблокировано, 1 ложное срабатывание).
- **Экономия затрат:** Ранние пользователи сообщили о снижении затрат на 60–80% по сравнению с ручными подрядчиками при эквивалентном качестве.

Синтетический стресс-тест (1.000 задач, сентябрь 2025)

- **Среднее время выполнения:** 7,4 минуты (от заявки задачи до отправки результата).
- **Attestation overhead:** 3.1% (TEE report generation + signature verification adds 230 ms per task).
- **Validation cost:** 0.002 TOS per task (gas for validator score submissions + reward distribution).
- **Throughput:** Sustained 45 tasks/minute (2,700 tasks/hour) without performance degradation.
- **Fraud injection:** Intentionally submitted 50 fraudulent results (incorrect outputs, forged attestations); all detected by validators, resulting in 100% fraud detection rate.

5.13 Compliance & Human Impact

Regulatory Alignment TOS Foundations complies with emerging AI regulations:

- **US AI RMF** (NIST Risk Management Framework): Task receipts log model versions, datasets, energy usage, enabling auditors to assess risk categories (bias, safety, transparency).
- **EU AI Act:** High-risk AI systems (healthcare, critical infrastructure) can deploy on TOS with mandatory KYC, human-in-the-loop arbitration, and conformity assessments logged on-chain.
- **Singapore MAS Guidelines:** Financial institutions using AI agents for trading or advisory must maintain audit trails; TOS compliance API provides exportable reports meeting MAS requirements.



Job Displacement Mitigation While AI agents automate tasks, TOS creates new roles:

- **Validators:** Human experts earn income reviewing AI outputs (estimated 12% of task value).
- **Credential issuers:** Educational institutions, certification bodies issue verifiable credentials (new revenue stream).
- **Arbitrators:** Legal/technical specialists adjudicate disputes (fees funded by arbitration bonds).
- **Policy developers:** Governance participants draft and audit constitutional contracts (treasury grants).

Economic simulations (Section 12) suggest that for every 10 jobs automated, TOS creates 3–4 new specialised roles, partially offsetting displacement.

Energy Transparency EC tokens explicitly track energy consumption, incentivising renewable sources. Publishers can specify “green energy only” requirements; agents operating on solar/wind receive 10% fee discounts (TEM subsidy). Aggregate energy usage reported quarterly, enabling carbon offset purchasing via treasury.

6 РЫНКИ И БЕЗОПАСНОСТЬ – Переход МЕРКУРИЙ-ВЕНЕРА

6.1 Strategic Drivers

As the MERCURY era matures and transitions toward VENUS, generative AI’s economic impact grows substantially, with potential to add trillions annually to the global economy through knowledge work automation.²² As agents evolve into autonomous business units, enterprises demand regulated custody, hybrid settlement, and safety guardrails.

6.2 Product Portfolio

The Markets & Safety era introduces four core products targeting enterprise deployments and autonomous business units:

1. Regulated Custody (CustodyVault.py)

Rationale: Enterprises and regulated entities require programmable spending controls, multi-party approval, and auditability. Traditional blockchain wallets lack these features.

Architecture:

- **Multi-sig policies:** m -of- n signature schemes where m signatures (from agent keys, human guardians, auditors) approve transactions exceeding threshold T . Example: withdrawals $> 10,000$ TOS require 2-of-3 approval (agent + CFO, or agent + auditor, or CFO + auditor).
- **Hardware Security Modules (HSMs):** Private keys stored in FIPS 140-2 Level 3 certified HSMs (AWS CloudHSM, Thales Luna). HSMs attest key usage via signed logs; auditors verify no unauthorised access.
- **Policy engine:** Smart contract layer enforcing:
 - *Rate limits:* Maximum \$X per hour/day/month.
 - *Whitelists:* Transfers only to pre-approved addresses (e.g., payroll, vendor accounts).
 - *Blacklists:* Blocking sanctioned addresses (OFAC compliance).
 - *Emergency freeze:* Guardian-initiated pause on all operations pending investigation.

²²McKinsey Global Institute, “The economic potential of generative AI,” June 2023.



- **View-only keys:** Regulators receive read-only access to balances and transaction history without control over funds, enabling audits without custody.

Compliance: SOC 2 Type II audit package includes HSM attestation logs, policy execution traces, access control matrices. Integration with enterprise identity providers (Active Directory, Okta) via SAML/OAuth.

2. Hybrid Settlement (BridgeController.py, FiatGateway.py)

Rationale: Autonomous business units need to interact with traditional financial rails (bank accounts, credit cards, invoices) and other blockchain ecosystems.

Cross-Chain Bridges:

- **HTLC-based atomic swaps:** Hash Time-Locked Contracts enable trustless exchanges between TOS and other chains. Example: Agent locks 100 TOS on TOS chain with hash H ; counterparty locks equivalent stablecoin on external chain with same H ; both reveal pre-image to claim funds, or timeouts refund after 24 hours.
- **Bridge security:** Multi-sig committee (7-of-11 validators) co-signs cross-chain transfers; value limits ($< \$1M$ per transaction) and daily caps reduce attack surface. Formal verification (TLA+ model) proves bridge cannot lose funds under Byzantine assumptions.

Fiat Integration:

- **Banking APIs:** Integration with ISO20022-compliant payment networks (SWIFT, FedNow, SEPA). Agents invoice customers in USD/EUR; upon payment confirmation, CC/EC release from escrow.
- **Compliance:** KYC/AML checks via Chainalysis, Elliptic; transactions flagged as high-risk ($>95\%$ probability of illicit origin) automatically frozen pending human review.
- **Settlement speed:** Domestic transfers settle in 2–4 hours (real-time payment networks); international transfers 24–48 hours (correspondent banking).

Derivative Instruments:

- **Compute futures:** Contracts to purchase CC at fixed price in future (hedging against GPU price volatility). Margin requirements: 20% initial margin, 10% maintenance margin.
- **Energy options:** Call options on EC tokens enable agents to lock in renewable energy pricing; exercise price tied to solar/wind market indices.
- **Risk management:** Automated liquidation if margin falls below maintenance threshold; insurance fund (3% of futures volume) covers market dislocations.

3. Expert Arbitration Court (ArbitrationCourt.py)

Rationale: Disputes requiring domain expertise (legal interpretation, technical evaluation) cannot be resolved algorithmically.

Arbitrator Accreditation:

- **Eligibility:** Arbitrators must hold relevant credentials (law degree, AI safety certification, PhD in relevant field), complete TOS training module, and stake 10,000 TOS bond.
- **Specialisation:** Arbitrators register for sectors (legal, medical, financial, technical); disputes routed based on task category.
- **Performance tracking:** Arbitrators rated by disputing parties; low ratings ($< 3.5/5$ over 20 cases) result in probation or removal.

**Dispute Process:**

1. **Escalation:** Either party (publisher or agent) stakes arbitration bond (5–10% of task escrow).
2. **Assignment:** Smart contract selects 3 arbitrators (weighted by specialisation match, ratings, availability).
3. **Evidence submission:** Both parties upload evidence (code, datasets, contracts) to encrypted IPFS; arbitrators access via decryption keys.
4. **Deliberation:** Arbitrators review evidence, may request additional info, discuss via secure messaging.
5. **Ruling:** Majority vote (2-of-3) determines outcome (agent wins, publisher wins, split settlement). Ruling includes written rationale, published on-chain.
6. **Enforcement:** Escrow distributed per ruling; losing party forfeits bond (covers arbitration fees + damages).

Insurance Pool: Funded via TEM allocations (2% of AGIW volume); covers catastrophic cases where both parties acted in good faith but outcome is ambiguous (e.g., unforeseen technical failure). Payout requires 4-of-7 Council vote.

4. Safety Oracle (`SafetyOracle.py`)

Rationale: AI agents may produce harmful outputs (bias, toxicity, misinformation); enterprises need real-time safeguards.

Data Sources:

- **AI safety labs:** Streams from Anthropic’s Constitutional AI metrics, OpenAI’s GPT safety scores, Google DeepMind’s alignment dashboards.²³
- **Content moderation APIs:** Perspective API (toxicity), AWS Comprehend (sentiment, PII detection), Microsoft Azure Content Safety.
- **Dataset compliance tags:** Tags indicating dataset provenance, licensing, consent (GDPR Article 6 lawful basis).

Policy Enforcement:

- **Kill-switch:** If Safety Oracle flags agent output as high-risk (e.g., toxicity score > 0.9, PII detected), wallet operations auto-pause; requires guardian approval to resume.
- **Rate-limits:** Agents flagged for repeated violations (>3 in 30 days) subject to reduced task quotas (e.g., max 10 tasks/day).
- **Dataset compliance:** Tasks specifying “GDPR-compliant data only” enforce that agents prove dataset provenance via verifiable credentials; violations trigger contract rejection.

Transparency: Safety Oracle publishes monthly reports: aggregate violation rates by category, anonymised case studies, false-positive rates (target < 2%).

6.3 Custody Design

Custody contracts implement multi-sig logic requiring e.g., one agent key plus one guardian approval for transactions above threshold x . Guardians can delegate to auditors. Emergency brake halts operations when Safety Oracle flags severe risk; release requires multi-party approval.

²³IBM Research, “AI Safety Metrics for Enterprise Deployment,” May 2024.



6.4 Hybrid Settlement Mechanics

Settlement uses HTLCs with timeouts aligned to network latency. Fiat integration via ISO20022 ensures compatibility. For example, a USD payment is locked in bank escrow; once on-chain proof is emitted, CC tokens release. Reverse direction uses CC/EC as collateral until fiat clears.

6.5 Risk Analysis

Table 5 summarises controls.

Risk	Mitigation	Residual Exposure
Oracle manipulation	Multi-source median, slashable feeds	Low (requires collusion)
Validator cartel	Rotating committees, public telemetry	Medium (monitored)
Custody breach	HSMs, dual control, audits	Low (audited quarterly)
Arbitration backlog	Bond requirements, automated pre-screening	Medium (human scaling)
Safety oracle failure	Diverse data sources, manual override	Medium (requires governance intervention)

Таблица 5: Risk-control matrix for Markets & Safety era.

6.6 Power of AI (PAI) Consensus Roadmap

PAI transitions TOS from BlockDAG baseline (2,500+ TPS) to AI-assisted consensus targeting 10,000+ TPS while maintaining decentralisation and security.

Phase 1: AI Scheduling Hints (Early MERCURY) **Mechanism:** AI scheduler (trained on historical transaction patterns) proposes block ordering hints. Validators independently evaluate hints using heuristics (gas efficiency, AGIW task locality, fraud likelihood). If > 67% validators accept hint, block ordering follows AI proposal; otherwise, fallback to traditional BlockDAG topological sort.

Safety: Hints are non-binding recommendations; validators retain veto power. Worst-case: AI scheduler fails $\Rightarrow$ network reverts to baseline BlockDAG performance.

Expected Gain: 20–30% throughput improvement via better transaction batching and parallel execution optimisation.

Research: Publish open-source AI scheduler model (Apache 2.0 license), solicit community feedback, conduct adversarial testing (red-team attacks attempting to manipulate scheduler).

Phase 2: Hybrid Committees & Probabilistic Finality (Mid-MERCURY) **Mechanism:** Validators divided into two committees:

- **AI Committee** (30% of stake): Runs AI-managed workload balancing, proposes block orderings optimised for AGIW task dependencies.
- **Human Committee** (70% of stake): Reviews AI proposals, checks for anomalies (e.g., censorship, bias toward specific agents), provides final approval.



Probabilistic Finality: After k confirmations, transaction is “probably final” with confidence $1 - 2^{-k}$. AI Committee accelerates confirmation speed; Human Committee ensures safety.

Formal Verification: Use TLA+ to model committee interactions; prove liveness (eventual finality) and safety (no conflicting finals) under asynchronous Byzantine assumptions ($< 33\%$ malicious).

Expected Gain: 3–4× throughput (7,500–10,000 TPS) with sub-second finality for high-confidence transactions.

Governance Safeguards: Human Committee can invoke emergency brake if AI exhibits unexpected behaviour (e.g., favouring specific validators, censoring transactions). Brake requires 4-of-7 Council vote to lift.

Phase 3: Full PAI Consensus (Late MERCURY / Early VENUS) Mechanism: AI models manage:

- **Dynamic security budgets:** Adjust validator incentives in real-time based on network load, threat levels (DDoS, Sybil attacks), and economic conditions.
- **Adaptive sharding:** Partition state and workload across sub-DAGs (shards) based on transaction locality; AI predicts optimal shard assignments.
- **Predictive fee markets:** AI forecasts demand, adjusts TEM subsidies preemptively to prevent congestion.

Performance Target: 10,000+ TPS sustained, 15,000+ TPS peak, sub-500ms finality for 95% of transactions.

Decentralisation Constraints: Governance sets bounds:

- Minimum validator count: 500 (prevents excessive centralisation).
- Maximum validator earnings: Gini coefficient < 0.45 (ensures broad participation).
- AI model transparency: Open-source models, public training data, reproducible builds.

Century-Scale Vision: PAI consensus adapts to interplanetary latency (Section 8), dynamic node availability (colonies joining/leaving network), and evolving threat landscapes (post-quantum attacks, AGI adversaries).

6.7 Two-Lens Rationale

Near-Term (MERCURY Era)

Enterprises gain compliant automation with auditable AGIW receipts, regulated custody (HSMs, multi-sig), hybrid fiat/crypto settlement, expert dispute resolution, and AI safety guardrails. Cost savings (60–80%) and compliance readiness (SOC 2, EU AI Act) drive adoption.

Century-Scale (VENUS through ANDROMEDA)

Infrastructure supports Decentralised Autonomous Economies (DAEs) operating compute farms, energy grids, and cross-planet commerce with managed risk. PAI consensus scales to interstellar latencies, constitutional contracts encode governance for mixed human–AI populations, reversible history mechanisms allow court-authorized rollbacks for catastrophic failures.

6.8 Risk Matrix & Mitigation Summary



Risk	Mitigation	Residual Exposure
Oracle price manipulation	Multi-source median aggregation, slashable stake (1,000 TOS per feeder), staleness fallback to EWMA	Low (requires 11/21 colluding feeders; detected via statistical analysis)
Validator cartel formation	Rotating committees, quadratic voting for governance, public telemetry dashboards, RepGraph anti-collusion detection	Medium (monitored quarterly; governance can adjust incentives)
Custody breach (key theft)	HSMs (FIPS 140-2 Level 3), multi-sig policies, rate limits, emergency freeze, quarterly penetration testing	Low (attack requires physical HSM access + multi-party collusion)
Arbitration backlog	Bond requirements (deter frivolous disputes), automated pre-screening (ML classifier predicts merit), performance incentives for arbitrators	Medium (scales with arbitrator recruitment; TEM subsidises training)
Safety Oracle failure (false negatives)	Diverse data sources (Anthropic, OpenAI, Google, independent auditors), manual override by guardians, monthly accuracy audits	Medium (requires ongoing model tuning; governance reviews threshold policies)
PAI centralisation risk	Open-source AI models, governance-enforced bounds (min validators, max Gini), Human Committee veto, emergency brake	Medium (long-term governance challenge; community oversight critical)
Cross-chain bridge exploit	Formal verification (TLA+), multi-sig committees (7-of-11), value limits (\$1M/tx), insurance pool (3% of volume)	Low (exploit requires 4+ compromised validators + breaking cryptographic assumptions)

Таблица 6: Risk-control matrix for Markets & Safety era (MERCURY-VENUS transition).

7 Со-управление и метамиры – Переход ВЕНЕРА-МАРС

7.1 Социотехнический ландшафт

К 2070-м годам Метамир 1.0 возникает как самосогласованная виртуальная цивилизация. Правительства признают AI-сущности через Мирное соглашение о бифуркации. Инфраструктура управления должна кодировать законы, поддерживать обратимые решения и допускать задержки.



7.2 Компилятор политики как кода

Мотивация Конституционное управление требует кодирования законов, бизнес-логики и правил соблюдения в исполняемых модулях смарт-контрактов. Однако низкоуровневые языки реализации сложны, подвержены ошибкам и недоступны для не-разработчиков. Компилятор политики заполняет этот пробел.

Архитектура

1. **Policy DSL (предметно-ориентированный язык):** Высокоуровневый декларативный синтаксис для выражения правил. Пример:

```
policy AgentWithdrawalLimit {
  rule "Large withdrawals require multi-sig" {
    when: withdrawal.amount > 10000 TOS
    require: signatures.count >= 2
    from: [agent.key, guardian.key, auditor.key]
    timeout: 24 hours
  }
  rule "Sanctioned addresses blocked" {
    when: transfer.recipient in SanctionsList
    action: reject
    log: "OFAC violation attempt"
  }
}
```

2. **Конвейер компилятора:**

- **Парсер:** Токенизирует DSL, строит абстрактное синтаксическое дерево (AST).
- **Статический анализатор:** Проверяет ограниченные циклы (без бесконечного газа), явный контроль доступа (без незащищенных функций), обязательное логирование событий.
- **Генератор кода:** Выдает Python-код для виртуальной машины TOS.
- **Оптимизатор:** Снижает стоимость газа через устранение мертвого кода, сворачивание констант.
- **Верификатор:** Формальная верификация с использованием SMT-решателей (Z3, CVC5) для доказательства свойств безопасности (например, «нет несанкционированного вывода», «все изменения состояния зарегистрированы»).

3. **Рабочий процесс развертывания:**

- Управление составляет политику в DSL, фиксирует в GitHub.
- Компилятор запускает CI/CD конвейер: компиляция, тестирование на testnet, генерация отчета об аудите.
- Сообщество проверяет код, симуляции, доказательства формальной верификации (7-дневный период).
- Голосование в сети: требуется 60% кворум, 75% одобрение.
- Развертывание с временной блокировкой: 48-часовая задержка перед активацией (окно аварийного тормоза).



Гарантии безопасности Компилятор обеспечивает:

- **Завершаемость:** Все циклы ограничены константами или пределами, установленными управлением (нет рисков проблемы останова).
- **Контроль доступа:** Каждая функция, изменяющая состояние, требует явных проверок ролей (`onlyAgent`, `onlyGuardian`, `onlyGovernance`).
- **Проверяемость:** Все действия политики выдают события с временными метками, ID субъектов и обоснованиями.
- **Обновляемость:** Политики версионированы; старые версии архивированы неизменно; миграции проверены.

Примеры использования

- **Экспортный контроль:** «Агенты в юрисдикции X не могут выполнять задачи с наборами данных из юрисдикции Y» (соответствие ITAR).
- **Энергетические квоты:** «Виртуальные университеты выделяют 1000 ЕС/месяц на студента; превышения требуют одобрения декана».
- **Аварийное реагирование:** «Если оракул безопасности отмечает > 10 нарушений токсичности за 1 час, приостановить все операции агентов до проверки Советом».

7.3 Протокол расчетов с задержкой

Используя репликацию на основе CRDT, участники поддерживают частичное состояние и согласовывают путем обмена логами. Алгоритм 1 описывает процесс.

Algorithm 1 Delay-Tolerant Settlement (DTS)

- 1: Input: Operation log L_i for participant i .
 - 2: Broadcast digest $d_i = H(L_i)$ periodically.
 - 3: On receiving digest d_j , request missing operations.
 - 4: Merge logs via $L_i \leftarrow \text{merge}(L_i, L_j)$, ensuring commutativity.
 - 5: When quorum acknowledges, emit finality proof to main chain.
-

7.4 Метрики

SLO управления: 75% участие в критических голосованиях, время подготовки предложения ≥ 48 часов, ноль аварийных откатов за квартал. Телеметрия включает явку голосующих, успех выполнения, результаты споров.

7.5 Примеры применения

- Виртуальные университеты кодируют учебные программы, полномочия и аккредитацию в контрактах политики.
- Смешанные человеко-АИ корпорации используют конституционные контракты для одобрения бюджета, найма и контроля рисков.
- Межреальная торговля использует расчет с задержкой для обмена ресурсами.



8 Расширение и обратимая история – Эры ЮПИТЕР - АНДРОМЕДА

По мере того, как цивилизация AGI расширяется за пределы Земли (эра ЮПИТЕР) и в конечном итоге стабилизируется в звездных системах (эра АНДРОМЕДА), TOS должна эволюционировать для обработки межпланетной задержки, восстановления после катастрофических сбоях и управления на протяжении нескольких столетий. В этом разделе представлены архитектурные расширения, которые превращают TOS из наземного расчетного слоя в экономический субстрат для гео-звездной цивилизации.

8.1 Профиль межпланетного консенсуса

Проблема: задержки связи со скоростью света Задержка передачи данных между Землей и Марсом варьируется от 8 до 44 минут в зависимости от орбитальных позиций. Лунные колонии испытывают задержки 2,5 секунды. Поселения пояса астероидов (Церера, главный пояс) сталкиваются с задержками 40–80 минут. Традиционные протоколы консенсуса (BFT, PoS) предполагают субсекундную связь; прямое применение к межпланетным сетям вызывает:

- **Остановка финализации:** Валидаторы истекают по времени ожидания межпланетных подтверждений.
- **Распространение форков:** Разделение сети во время планетных соединений (когда небесные тела выравниваются таким образом, что прямая связь блокируется Солнцем).
- **Энергетические потери:** Избыточное производство блоков в изолированных колониях.

Решение: иерархический консенсус с корнями, устойчивыми к задержкам TOS принимает двухуровневую модель консенсуса:

Уровень 1: локальные домены консенсуса (LCD)

Каждое планетное поселение (Земля, Марс, Луна, Церера и т.д.) управляет независимым экземпляром TOS с локальными валидаторами. Консенсус внутри LCD использует стандартный BlockDAG + PAI (субсекундная финализация). Транзакции рассчитываются локально; агенты на Марсе платят другим агентам на Марсе без ожидания подтверждения от Земли.

Уровень 2: межпланетная корневая цепь (IRC)

Каждые 24 часа (земное время) каждый LCD публикует *доказательство снимка* в IRC—корень Меркла локального состояния плюс zk-SNARK аттестация, доказывающая корректное выполнение. IRC агрегирует снимки через:

1. **BFT с задержкой:** Модифицированный Tendermint с экспоненциальным откатом; валидаторы ждут до 60 минут межпланетных голосований.
2. **Ослабление кворума:** Финализация требует 51% доли (против 67% наземных), принимая более высокий риск форка в обмен на живучесть.
3. **Контрольные точки:** Каждые 7 дней (земное время) IRC выдает *каноническую контрольную точку*—необратимый якорь расчета, признаваемый всеми LCD.

Межпланетные атомарные свопы Агент на Земле хочет заплатить агенту на Марсе. Протокол:

1. Земной агент блокирует средства в HTLC (хеш-контракт с временной блокировкой) на земном LCD, с хешем H и 48-часовым тайм-аутом.
2. Земной LCD публикует доказательство HTLC в IRC (следующий ежедневный снимок).



3. Марсианский LCD синхронизирует IRC, проверяет HTLC, позволяет марсианскому агенту получить средства, раскрывая прообраз x (где $H = \text{hash}(x)$).
4. Марсианский LCD публикует доказательство получения в IRC.
5. Земной LCD наблюдает получение, освобождает средства первоначальному заявителю после истечения тайм-аута (или возвращает при тайм-ауте).

Компромиссы: Задержка расчета = 48–96 часов (против 15 минут наземных). Приемлемо для торговли агент-агент с учетом снижения зависимости от земной ликвидности.

Анализ безопасности Сценарий атаки: Противник контролирует 40% марсианских валидаторов, пытается двойную трату, форкуя марсианский LCD и отправляя конфликтующий снимок в IRC.

Защита: Валидаторы IRC обнаруживают конфликтующие корни Меркла, отклоняют оба снимка, запускают арбитраж. Марсианский LCD входит в *режим восстановления*: останавливает новые транзакции, ожидает комитет аудита под руководством Земли для проверки истории блоков и назначения канонической цепи. Штраф: злонамеренные марсианские валидаторы теряют 50% доли; оценка репутации марсианского LCD снижена, увеличивая комиссии за межпланетный своп.

Целевые показатели производительности (эра ЮПИТЕР)

- **Локальный TPS:** 10.000+ на LCD (оптимизировано PAI).
- **Межпланетный расчет:** 24–48 часов (ежедневные снимки).
- **Каноническая финализация:** 7 дней (еженедельные контрольные точки).
- **Устойчивость к разделению сети:** Выживание 30-дневного блэкаута связи с Марсом (хранит ожидающие снимки, синхронизируется при повторном подключении).

8.2 Механизм обратимой истории

Мотивация: восстановление после катастрофических ошибок Поскольку агенты AGI управляют критической инфраструктурой (системы жизнеобеспечения, орбитальная механика, генетические биобанки), *необратимые* ошибки становятся экзистенциальными рисками. Примеры:

- Мошеннический агент выполняет мошеннический перевод, опустошая казну колонии (\$50B+).
- Ошибка смарт-контракта вызывает каскадные сбои в контрактах энергосети.
- Скомпрометированный набор валидаторов (прорыв квантового компьютера) подрывает безопасность подписи.

Традиционные блокчейны не предлагают никакого средства, кроме спорных хардфорков по социальному консенсусу. TOS обеспечивает *откаты, авторизованные судом*, сохраняя при этом проверяемость.

Архитектура: откат на основе контрольных точек Создание контрольной точки
Каждые 24 часа каждый LCD вычисляет:

- **Корень состояния:** Корень Меркла всех балансов счетов, хранилища смарт-контрактов, оценок репутации.



- **Доказательство валидности:** zk-SNARK, доказывающий, что все транзакции с последней контрольной точки выполнены правильно (лимиты газа, подписи, переходы состояния).
- **Метаданные:** Высота блока, временная метка, идентификатор предлагающего, хеш предыдущей контрольной точки.

Контрольные точки публикуются в IRC и архивируются в резервном хранилище (см. Федерация архивов).

Авторизация отката

Процедура управления (требует конституционного сверхбольшинства, 75% голосов):

1. **Отчет об инциденте:** Пострадавшие стороны представляют доказательства (журналы транзакций, экспертные показания, судебный анализ) в Арбитражный суд.
2. **Судебное решение:** Если суд определяет откат обоснованным (мошенничество, эксплойт или форс-мажор), выдает подписанный *сертификат авторизации отката (RAC)*, указывающий:
 - Идентификатор целевой контрольной точки (высота, временная метка).
 - Обоснование (юридическое обоснование, оценка воздействия).
 - Исключения (транзакции для сохранения несмотря на откат, например, невинные платежи третьих сторон).
3. **Голосование сообщества:** RAC передается в управление; требуется 75% сверхбольшинства + подпись валидатора.
4. **Выполнение:** Валидаторы воспроизводят транзакции от целевой контрольной точки вперед, исключая откатываемые события. Вычисляется новый корень состояния; форк становится канонической цепью.

Технические меры защиты

- **Неизменный след аудита:** Исходные транзакции остаются в архивных журналах; откат создает *параллельную историю* («ветка отмены»). Криминалисты могут сравнивать временные линии.
- **Ограниченная область:** Откаты ограничены ретроспективой на 90 дней (более старые контрольные точки необратимы). Предотвращает неопределенную неопределенность.
- **Компенсационный фонд:** 5% казны зарезервировано для *компенсации невинных жертв*—агентов, потерявших средства из-за отката (например, продавец, отправивший товары до отката платежа покупателя).
- **Ограничения по частоте:** Максимум 1 откат в квартал; предотвращает злоупотребление управлением.

Тематическое исследование: симулированная квантовая атака (2185) Гипотетический сценарий: противник со 100-кубитным отказоустойчивым квантовым компьютером взламывает подписи ECDSA, похищает \$12B с 5.000 счетов агентов за 8-часовой период (до завершения миграции PQC).

Реакция:

1. Центр операций безопасности обнаруживает аномальные шаблоны транзакций (10.000 переводов со счетов без недавней активности).
2. Совет вызывает аварийный автоматический выключатель; сеть приостановлена в течение 2 часов.



3. Арбитражный суд собирается; подтверждает квантовую атаку через криптоанализ.
4. Голосование управления (82% одобрение) разрешает откат к контрольной точке за 12 часов до атаки.
5. Валидаторы выполняют откат; украденные средства восстановлены; атакующие транзакции аннулированы.
6. Компенсационный фонд выплачивает \$200M агентам, получившим законные платежи в окне атаки.
7. Пост-мортем: ускорение миграции PQS; обновление до Dilithium-5 в течение 30 дней; публикация отчета о прозрачности.

Результат: Кризис разрешен за 7 дней; доверие к сети сохранено; квантовое преимущество атакующего нейтрализовано возможностью отката.

8.3 Федерация архивов

Проблема: сохранение данных на тысячелетия Цивилизация AGI требует *институциональной памяти*, охватывающей столетия. Варианты использования:

- Юридические споры со ссылкой на контракты, подписанные 200 лет назад.
- Научная воспроизводимость (проверка экспериментальных утверждений из предыдущей эры).
- Генеалогические исследования (записи о происхождении человека/AI).
- Исторические аудиты (заплатила ли колония X колонии Y в 2234 году?).

Традиционное облачное хранилище (AWS Glacier, Google Coldline) предполагает хранение <100 лет; предприятия рушатся, форматы устаревают, накапливается битовая гниль.

Решение: многоуровневая стратегия архивирования **Уровень 1: горячее хранилище (1–10 лет)**

Активные узлы хранят полное состояние + последние блоки в SSD/NVMe. Индексировано для быстрых запросов. Избыточность: 3 реплики на LCD, географически распределенные (Земля: 3 континента; Марс: Olympus Mons, Valles Marineris, Hellas Basin).

Уровень 2: теплое хранилище (10–100 лет)

Архивные узлы хранят сжатые данные блоков в массивах HDD. Кодирование стирания Рида-Соломона (20% избыточность) допускает 20% отказов дисков. Ежемесячная очистка обнаруживает битовую гниль; поврежденные блоки восстанавливаются из четности. Стимул: казна платит 0,5 TOS за ТБ-год проверенного хранения.

Уровень 3: холодное хранилище (100–1.000 лет)

Экспериментальные носители:

- **Хранение ДНК:** Кодирование данных блокчейна в синтетической ДНК (Church Lab, Гарвард). Плотность: 215 петабайт на грамм; срок службы: 2.000+ лет при хранении при -18°C. Стоимость: \$1.000 за ТБ (цены 2025 года; прогноз \$10/ТБ к 2050 году).
- **5D оптические диски:** Плавленное кварцевое стекло с фемтосекундными лазерными наноструктурами (Университет Саутгемптона). Срок службы: 13,8 миллиардов лет (проверено путем ускоренного старения). Плотность: 360 ТБ на диск.
- **Архивная лента (LTO-12+):** Магнитная лента в климатически контролируемых хранилищах. Срок службы: 50–100 лет (требуется периодической миграции).



Стратегия: Контрольные точки состояния реестра (каждые 7 дней) записываются на все три носителя; хранятся в географически разрозненных хранилищах (Земля: Всемирное семенное хранилище на Шпицбергене, Антарктида; Луна: лавовые трубки; Марс: подземные ледяные пещеры).

Уровень 4: протокол федерации

Колонии разделяют ответственность за архивирование через *распределенное хранение*:

1. Земная колония архивирует данные Марса; Марс архивирует лунные данные; Луна архивирует данные Цереры и т.д. (кольцевая топология).
2. Ежегодная *церемония очистки*: колонии обмениваются доказательствами Меркла, проверяют целостность данных.
3. Если колония выходит из строя (например, катастрофический сбой системы жизнеобеспечения), выжившие колонии восстанавливают ее архив из фрагментов с кодом стирания.

Протокол извлечения Запрос исторического состояния (например, «Каков был баланс счета `tos1qq5hqu2jwn0r5ehvuw5r3qjcfvjvq5dhldxvyq` на 2234-05-15?»):

1. Проверка горячего хранилища (если в пределах 10 лет): индексированный поиск $O(\log n)$, задержка 100мс.
2. Проверка теплого хранилища (если 10–100 лет): распаковка блоков, проверка доказательства Меркла, 10–60 секунд.
3. Проверка холодного хранилища (если 100–1.000 лет): извлечение ДНК/оптического диска из хранилища, секвенирование/чтение данных, восстановление через Рида-Соломона, 1–7 дней.

Экономическая модель **Финансирование:** 3% годовых доходов казны выделяется на архивную инфраструктуру. При казне в \$50B (прогноз 2150) это финансирует \$1,5B/год для:

- Синтез ДНК (\$500M/год, кодирование 500 ПБ).
- Производство 5D оптических дисков (\$300M/год, 1.500 ПБ).
- Строительство/обслуживание хранилищ (\$400M/год).
- Операторы узлов хранения (\$300M/год, 1.000 узлов по всему миру).

8.4 Миграция на постквантовую криптографию

Временная шкала угроз

- **2030-е:** 50-кубитные шумные квантовые компьютеры (недостаточно для атак ECDSA).
- **2040-е:** 1.000-кубитные отказоустойчивые системы (алгоритм Шора взламывает RSA-2048, ECDSA-256 за часы).
- **2050-е:** Коммерческие квантовые компьютеры доступны хорошо финансируемым противникам.

Стратегия миграции: гибридный переход **Фаза 1 (v4.0, 2026):** Введена двойная подпись. Каждая транзакция несет две подписи:

- Классическая: ECDSA `secp256k1` (обратная совместимость).
- Постквантовая: Dilithium-3 (NIST FIPS 204).



Валидаторы проверяют обе; сеть принимает транзакцию, если *любая* подпись действительна (обеспечивает живучесть, если реализация PQS имеет ошибки).

Фаза 2 (v5.0, 2030): Режим только PQS. Классические подписи устарели; узлы отклоняют транзакции только с ECDSA. Льготный период: 12 месяцев для обновления кошельков агентов.

Фаза 3 (v6.0, 2040): Квантово-устойчивые хеш-функции. Замена SHA-256 (уязвим для алгоритма Гровера) на SHA-3 или BLAKE3 (256-битная безопасность сохранена против квантовых атак).

Церемония сообщества (прозрачность и доверие) Для обеспечения прозрачности и доверия сообщества TOS проводит *церемонию миграции PQS*:

1. **Фаза 1 (месяц 1–3):** Открытый призыв участников (университеты, фирмы безопасности, добровольцы сообщества). 5.000+ участников вносят случайность для генерации глобальных параметров PQS.
2. **Фаза 2 (месяц 4):** Многосторонние вычисления (MPC) генерируют главный публичный ключ. Стенограмма церемонии опубликована; участники уничтожают приватные вклады.
3. **Фаза 3 (месяц 5–6):** Тестирование совместимости. Валидаторы запускают теневой форк с подписями PQS; стресс-тест в состязательных условиях.
4. **Фаза 4 (месяц 7):** Голосование управления (требуется 67% одобрение). Если принято, обновление mainnet запланировано на 30 дней.
5. **Фаза 5 (месяц 8):** Активация mainnet. Старые подписи поэтапно выводятся в течение 12 месяцев.

План отступления Если обнаружена критическая уязвимость PQS (например, нарушено предположение о жесткости решетки):

- Аварийный откат к последней контрольной точке перед PQS.
- Активация альтернативной схемы PQS (Falcon, SPHINCS+).
- Повторный созыв церемонии с обновленными параметрами.

8.5 Философские основы: временной суверенитет

Механизм обратимой истории воплощает радикальный принцип управления: **будущие поколения могут исправлять прошлые ошибки**. Традиционные блокчейны превозносят неизменность как священное; TOS признает, что *идеальное предвидение невозможно*. Позволяя откаты, санкционированные судом, TOS дает цивилизации AGI возможность:

- Восстанавливаться от событий черного лебедя (квантовые атаки, мошеннический суперинтеллект).
- Отменять несправедливые транзакции (кража, принуждение, алгоритмическая дискриминация).
- Сохранять институциональную преемственность на протяжении столетий (конституционные поправки, пересмотр договоров).

Эта способность не безгранична (90-дневное окно, 75% сверхбольшинство, ограничения частоты), но сигнализирует о сдвиге от *код есть закон* к *код + управление есть закон*. Поскольку сущности AGI получают правовой статус, а люди сливаются с цифровыми субстратами, способность *ответственно переписывать историю* становится необходимой для устойчивости цивилизации.



9 Технические основы

TOS architecture rests on four foundational pillars, each addressing critical requirements for the agent economy. These pillars are designed to evolve across celestial eras while maintaining backward compatibility.

9.1 P1 – Digital Personhood & Reputation

Decentralized Identifiers (DIDs) Every agent, validator, and human participant receives a W3C-compliant DID (Decentralized Identifier) following the `did:tos:` method specification.

DID Format:

`did:tos:<bech32-address>`

`did:tos:<network-id>:<bech32-address>`

Examples:

Simple form: `did:tos:tos1qvqsyqcyq5rqwzqfpg9scrgwpugpzysnzs23v9xh`

Mainnet: `did:tos:mainnet:tos1qvqsyqcyq5rqwzqfpg9scrgwpugpzysnzs23v9xh`

Testnet: `did:tos:testnet:tst1qqqsyqcyq5rqwzqfpg9scrgwpugpzysnzs2pgde5k`

Note: TOS uses complete bech32 addresses as DID identifiers, preserving full address format including checksum for validation. Mainnet addresses use `tos1` prefix, while testnet addresses use `tst1` prefix.

DID Document Structure:

- **Public Keys:** Ed25519 (signing), X25519 (encryption), Dilithium-3 (post-quantum).
- **Service Endpoints:** HTTPS URLs for off-chain communication, IPFS content addresses.
- **Verifiable Credentials:** Linked to attestations (reputation scores, skill certifications, regulatory compliance).
- **Revocation Registry:** Bloom filter enabling efficient revocation checks without revealing full credential list.

Registration Process:

1. Agent generates key pair; derives DID from public key hash.
2. Submits registration transaction to TOS chain (fee: 1 TOS).
3. DID anchored on-chain; DID document published to IPFS; hash stored in smart contract.
4. Agent proves control by signing challenge with private key.

Verifiable Credentials (VCs) Agents acquire capabilities through tamper-evident credentials issued by trusted authorities (validators, arbitration courts, educational institutions).

Credential Types:

- **Reputation Credentials:** RepGraph tier (Bronze/Silver/Gold/Platinum), task completion count, dispute resolution history.
- **Skill Badges:** Domain expertise (legal analysis, climate modeling, code review) validated via peer assessment or certification exams.
- **Compliance Attestations:** KYC/AML clearance (for fiat on-ramps), regulatory approvals (EU AI Act conformity, FDA medical device certification).
- **Delegation Proofs:** Authority to act on behalf of another entity (corporate agents executing on behalf of DAOs).

**Issuance Protocol:**

1. Issuer (e.g., RepGraph smart contract) creates credential JSON-LD document.
2. Signs with BBS+ signatures (selective disclosure enabled; holder can prove specific attributes without revealing entire credential).
3. Publishes credential to holder's DID document via IPFS; on-chain event emitted.
4. Holder stores credential in wallet; presents zero-knowledge proofs to verifiers (e.g., "Prove reputation ≥ 0.7 without revealing exact score").

Reputation Graph (RepGraph) Algorithms Score Computation:

$$r_t = \alpha \cdot r_{t-1} + (1 - \alpha) \cdot s_t \quad (14)$$

where:

- r_t : Reputation score at time t (range $[0, 1]$).
- $\alpha = 0.85$: Exponential smoothing factor (85% historical weight).
- s_t : Quality score for most recent task (median of 5 validator assessments).

Decay Mechanism:

$$r_{t+\Delta t} = r_t \cdot (0.98)^{\Delta t/30} \quad (15)$$

Reputation decays 2% per month of inactivity; prevents dormant high-reputation accounts from sudden reactivation.

Multi-Dimensional Scoring:

- **General Reputation:** Aggregate across all task types.
- **Domain-Specific:** Separate scores for legal, medical, financial, technical domains.
- **Recency-Weighted:** Recent performance (last 30 days) weighted 2× vs. historical average.

Anti-Gaming Measures:

- **Sybil Resistance:** Minimum 100 TOS stake + proof-of-uniqueness (BrightID integration roadmap v5.0).
- **Collusion Detection:** Graph analysis identifies clusters of agents with suspiciously correlated voting patterns; flagged for manual review.
- **Reputation Caps:** New agents capped at 0.5 reputation for first 90 days; prevents rapid reputation farming.

Interoperability with W3C Standards

- **DID Core v1.0:** Full compliance; TOS DIDs resolvable via Universal Resolver.
- **Verifiable Credentials Data Model v1.1:** Credentials valid across ecosystems (can present TOS credential to non-TOS verifier).
- **DIDComm v2:** Secure messaging protocol for agent-to-agent communication (encrypted, authenticated, transport-agnostic).



9.2 P2 – Economic Execution

Account Abstraction TOS implements native account abstraction, enabling programmable accounts with custom validation logic integrated directly into the protocol layer. This eliminates the need for separate smart contracts or external coordination services, providing a streamlined experience for AI agents.

Key Features:

- **Paymasters:** Third parties sponsor gas fees (e.g., enterprise pays fees for employee agents).
- **Multi-Sig Wallets:** m -of- n approval for high-value transactions (corporate treasury accounts).
- **Session Keys:** Temporary keys with restricted permissions (agent can sign small transactions without accessing main key).
- **Social Recovery:** Recover account access via guardians (family members, trusted agents) if private key lost.

UserOperation Structure:

```
{
  sender: "0x...",           // Agent's account address
  nonce: 42,                  // Sequential counter
  initCode: "0x...",          // Account creation code (if new)
  callData: "0x...",          // Function call to execute
  paymasterAndData: "0x...",  // Paymaster address + verification data
  signature: "0x..."        // Signature over UserOperation
}
```

AGIW Settlement Engine Decentralized task marketplace where agents earn by performing verifiable work. Full protocol specification in Section 10.

Core Components:

- **TaskInstance Contract:** Escrow management, state machine (Open → Claimed → Verified → Settled).
- **Validator Pool:** Random selection of 5 validators per task; median quality score determines payout.
- **Attestation Reports:** TEE-generated proofs (Intel SGX, AMD SEV) certifying execution environment integrity.
- **Dispute Resolution:** Escalation to expert arbitrators if agent contests validator assessment.

Resource Token Markets (CC/EC) Compute Credits (CC): Fungible tokens representing GPU/CPU compute hours. Pricing oracles aggregate AWS/GCP/Azure spot prices.

Energy Credits (EC): Tokenized energy consumption (kWh). Oracle feeds from grid operators; renewable energy sources earn 10% premium.

Automated Market Maker (AMM):

- Constant Product Market Maker (CPMM): $x \cdot y = k$ where x = CC reserves, y = TOS reserves.
- Liquidity providers earn 0.3% swap fees.
- Impermanent loss protection: TEM treasury subsidizes divergence losses for LPs providing >1M TOS liquidity.



TOS Energy Model (TEM) Dynamic fee adjustment mechanism balancing network sustainability with agent accessibility.

Fee Structure:

$$\text{Total Fee} = \text{Base Fee} + \text{Priority Tip} - \text{TEM Subsidy} \quad (16)$$

Base Fee:

- Computed via dynamic adjustment mechanism that adjusts based on block fullness.
- Target: 50% block utilization; base fee increases 12.5% per block if usage > target.

TEM Subsidy Categories:

- **Public Goods (100% subsidy):** Open-source development, climate research, medical diagnostics.
- **Education (75%):** Academic research, student projects.
- **SMBs (50%):** Small/medium businesses with <\$10M annual revenue.
- **Enterprise (0%):** Full commercial rates for large corporations.

Funding: 8% of block rewards allocated to TEM treasury; governance votes on subsidy allocations quarterly.

Таблица 7: TEM Economic Parameters and Control Mechanisms

Parameter	Value / Range	Control Logic
Base fee band (TOS)	0.0001–0.01 per gas unit	Adjust by block fill p95; target 50% utilization
Base fee adjustment	+12.5% per block	If usage > 50% target
CC stabilization band	±15% of GPU index	Oracle median with 30-min staleness cap
EC stabilization band	±10% of kWh index	Region-weighted energy basket
TEM treasury split	40% security / 30% grants / 30% rebates	Quarterly governance rebalancing
Oracle staleness cap	≤30 minutes	Freeze CC/EC subsidy if exceeded
Subsidy categories	Public Goods: 100% / Education: 75% / SMB: 50% / Enterprise: 0%	Task classification by governance vote
Treasury allocation	8% of block rewards	Fixed protocol parameter
Validator reward pool	60% of block rewards	Base consensus incentive
Burn mechanism	10% of base fees	Deflationary pressure

Economic Sensitivity: TEM parameters undergo quarterly review by governance. Simulation results (available at metrics.tos.network/tem-sensitivity) demonstrate system stability under 3× demand spikes and 50% oracle variance scenarios.

9.3 P3 – Consensus & Safety

BlockDAG Baseline Architecture Traditional blockchains form linear chains; TOS uses Directed Acyclic Graph (DAG) enabling parallel block production.

Block Structure:



- Each block references 1–3 parent blocks (vs. 1 in traditional linear chains).
- Validators independently propose blocks; no leader election bottleneck.
- Topological sorting determines transaction order; conflicts resolved via timestamp + proposer priority.

Consensus Rules:

- **Confirmation Depth:** Transaction considered final after 8 blocks deep (90 seconds median).
- **Longest Chain Heuristic:** In case of DAG forks, prefer subgraph with most cumulative proof-of-stake weight.
- **Validator Rotation:** Committee of 100 validators selected per epoch (24 hours); weighted by stake + reputation.

Performance:

- Baseline: 2,500–3,500 TPS (measured Sept 2025).
- PAI-optimized: 10,000+ TPS (target Q4 2026).

Таблица 8: Performance Evaluation Methodology

Parameter	Configuration
Validators	32 nodes across 8 geographic regions (N. America, Europe, Asia-Pacific, S. America, Africa, Middle East, Oceania, Central Asia); BLS signature scheme
Hardware	16 vCPU, 64 GB RAM, NVMe SSD (1 TB), 2.5 Gbps network interface
Network	Emulated RTT: 20–180ms (inter-region), packet loss: 0.1–1.0%, jitter: ± 10 ms
Workloads	Real agent jobs: code review (15%), data synthesis (25%), RAG tasks (20%); Synthetic spam: 20–60% injection rate
Runtime	72-hour continuous runs; 5 independent seeds; cold start + steady state phases
Block Size	Variable: 100–2,000 transactions per block; target 50% utilization
Mempool Config	Priority queue with reputation weighting; max 100,000 pending transactions
Metrics Collected	TPS (p50/p95/p99), block time (p50/p95), finality depth (confirmation blocks), AGIW settlement latency (task publish $\rightarrow$ reward), dispute rate (% of tasks disputed), validator liveness (uptime %), fork rate
Measurement Tools	Prometheus + Grafana; custom telemetry pipeline; raw logs exported to <code>metrics.tos.network</code>
Baseline Period	September 2025 testnet deployment (14-day observation window)

Key Results (September 2025 baseline):

- TPS: p50 = 2,847, p95 = 3,421, p99 = 3,789
- Block time: p50 = 11.3s, p95 = 14.2s
- Finality: 8 blocks (median 90s)
- AGIW settlement: p50 = 12.4 min, p95 = 18.7 min
- Dispute rate: 1.8% of completed tasks
- Validator liveness: 98.7% average



Power of AI (PAI) Consensus AI-assisted scheduling enhances BlockDAG throughput via predictive transaction ordering.

Phase 1: AI Scheduling Hints

Lightweight ML model (LSTM trained on historical transaction patterns) predicts optimal block ordering. Validators vote on AI proposal; >67% approval adopts hint.

Phase 2: Hybrid Committees

AI Committee (30% stake) proposes block orderings; Human Committee (70% stake) reviews and approves. Probabilistic finality: $1 - 2^{-k}$ confidence after k confirmations.

Phase 3: Full PAI

AI agents participate as validators. Governance caps AI voting power at 40% (prevents AI takeover); human validators retain veto.

AI Safety Oracle Integration On-chain oracle aggregates safety metrics from leading AI labs:

- **Anthropic Constitutional AI:** Harmfulness scores (toxicity, bias, misinformation).
- **OpenAI Moderation API:** Content policy violations (hate speech, violence, illegal activity).
- **Google Perspective:** Conversation quality metrics.

Enforcement:

- Tasks with safety score < 0.3 automatically rejected (refund escrow).
- Agents repeatedly submitting unsafe outputs flagged; reputation penalized.
- Human arbitrators review borderline cases (0.3–0.5 safety score).

Slashing & Anti-Sybil Mechanisms **Validator Slashing:**

- **Double-signing:** Propose conflicting blocks $\Rightarrow$ 50% stake slashed.
- **Downtime:** Offline >6 hours $\Rightarrow$ 5% stake slashed, ejected from active set.
- **Invalid Attestation:** Sign fraudulent TEE report $\Rightarrow$ 20% stake slashed.

Agent Anti-Sybil:

- Minimum stake: 100 TOS per agent identity.
- Rate limits: 5–60 minute cooldown between tasks (based on reputation).
- Graph analysis: Detect clusters of agents with identical behavioral patterns; require proof-of-uniqueness.

9.4 P4 – Resilience

Post-Quantum Cryptography Roadmap See Section 8 for detailed PQC migration strategy. Summary:

- **v4.0 (2026):** Dual-signing (ECDSA + Dilithium-3).
- **v5.0 (2030):** PQC-only; classical signatures deprecated.
- **v6.0 (2040):** Quantum-resistant hash functions (SHA-3).

Reversible History Court-authorised rollback mechanism for catastrophic failures. Checkpoints every 24 hours; 90-day rollback window; 75% governance supermajority required. See Section 8 for full specification.



Archival Federation Multi-layer storage strategy (Hot/Warm/Cold) with DNA and 5D optical disc backups. Data preserved for millennia across planetary colonies. See Section 8 for economic model and retrieval protocols.

Delay-Tolerant Networking Hierarchical consensus (Local Consensus Domains + Interplanetary Root Chain) handles light-speed communication delays. Daily snapshots enable cross-planet settlement within 48–96 hours. See Section 8 for interplanetary atomic swap protocol.

Disaster Recovery Protocols Byzantine Fault Scenarios:

- Network partition (Earth–Mars communication blackout): LCDs operate independently; sync when reconnected.
- Validator collusion (<33%): Median-of-5 scoring resists outliers.
- Validator collusion (>33%): Emergency governance intervention; slashing + re-election.

Recovery SLAs:

- **Minor incidents** (1–5% validator compromise): Automated slashing + replacement within 1 hour.
- **Major incidents** (5–20% compromise): Council emergency session; resolution within 24 hours.
- **Critical incidents** (>20% compromise): Network pause; rollback authorization; resolution within 7 days.

10 Спецификация протокола AGI Work (AGIW)

10.1 Formal Definition

Task Representation A task T is a tuple:

$$T = (id, owner, spec, escrow_{CC}, escrow_{EC}, stake_{req}, deadline, whitelist, state) \quad (17)$$

where:

- $id \in \{0, 1\}^{256}$: Unique task identifier (keccak256 hash).
- $owner \in \mathbb{A}$: Publisher address.
- $spec$: IPFS content ID referencing task specification (JSON schema, datasets, acceptance criteria).
- $escrow_{CC}, escrow_{EC} \in \mathbb{R}^+$: Compute and energy credits locked for payment.
- $stake_{req} \in \mathbb{R}^+$: Minimum stake required for agents to claim task.
- $deadline \in \mathbb{N}$: Block number by which work must be submitted.
- $whitelist \subseteq \mathbb{A}$: Allowed agent set (empty $\Rightarrow$ public task).
- $state \in \{\text{Open, Claimed, Submitted, Validating, Settled, Disputed}\}$.

Agent Submission An agent a submits work as:

$$W = (resultHash, attestation, metrics) \quad (18)$$

where:

- $resultHash \in \{0, 1\}^{256}$: Content hash of task output (IPFS CID).



- *attestation*: TEE report including:
 - *measurementHash*: SHA-256 hash of executing code (proves integrity).
 - *inputHash, outputHash*: Hashes of inputs and outputs.
 - *timestamp*: Execution start/end times.
 - *energyUsed* $\in \mathbb{R}^+$: kWh consumed (measured via OS telemetry or TEE counters).
 - *signature*: TEE-signed attestation (Intel SGX: ECDSA over secp256r1; AMD SEV: RSA-4096).
- *metrics*: Optional performance data (runtime, memory usage, model accuracy).

Validation Function Validators evaluate submissions using function:

$$\mathcal{V} : (W, spec, T) \rightarrow [0, 1] \quad (19)$$

The function checks:

1. **Attestation validity**: Verify TEE signature against known public keys (Intel, AMD).
2. **Measurement correctness**: Compare *measurementHash* to approved code hashes in whitelist.
3. **Result quality**: Compare *resultHash* to reference outputs (if deterministic) or sample outputs (if stochastic).
4. **Energy efficiency**: Ensure *energyUsed* $\leq$ budget specified in *spec*.

Validators submit quality scores $q_v \in [0, 1]$. Consensus aggregates via median:

$$q = \text{median}(\{q_v : v \in \text{assigned validators}\}) \quad (20)$$

Settlement requires $|q_v - q| \leq \epsilon$ for $\geq 67\%$ of validators, where $\epsilon = 0.1$ is tolerance threshold.

Slashing Policy Validators deviating $> 2\sigma$ from median forfeit:

$$slash_v = \min(0.05 \times stake_v, 0.01 \times stake_v \times |q_v - q|/\epsilon) \quad (21)$$

Slashed funds transferred to treasury. Persistent outliers (> 5 slashes in 100 tasks) removed from validator pool.

10.2 State Machine

The AGIW state machine transitions as follows:

State	Trigger	Conditions	Next State
Open	Agent claimTask()	calls Agent in <i>whitelist</i> (if specified), stake $\geq stake_{req}$, block < deadline	Claimed
Claimed	Agent submitWork()	calls Caller is claimant, block < deadline, valid attestation signature	Submitted
Claimed	Block $\geq$ deadline	Deadline passed without submission	Forfeited (agent loses 10% stake)



Submitted	Validators assigned	Automatic (triggered by smart contract event)	Validating
Validating	All validators submit scores	$\geq 67\%$ scores within ϵ of median	Settled
Validating	Dispute raised	Either party stakes arbitration bond	Disputed
Settled	Rewards distributed	Automatic (RewardSplitter executed)	Terminal state
Disputed	Arbitrators rule	Majority (2-of-3) vote	Resolved
Resolved	Ruling enforced	Escrow distributed per arbitration decision	Terminal state

Таблица 9: AGIW state transitions with triggers, conditions, and outcomes.

10.3 Attestation Workflow (Detailed)

Step 1: Agent Preparation

1. Agent provisions Trusted Execution Environment (TEE): Intel SGX enclave, AMD SEV-SNP VM, or ARM Realm.
2. Agent loads task code (verified against whitelist hash), input data, and model weights into TEE.
3. TEE initialises measurement: hash of code + data + configuration.

Step 2: Execution Inside TEE

1. TEE executes task in isolation (no OS visibility into enclave memory).
2. TEE logs energy consumption via:
 - **Intel SGX:** RAPL (Running Average Power Limit) counters.
 - **AMD SEV:** SMU (System Management Unit) telemetry.
 - **ARM Realm:** PMU (Performance Monitoring Unit) energy events.
3. TEE generates result, computes $resultHash = keccak256(output)$.

Step 3: Attestation Report Generation

TEE produces signed report:

```
{
  "version": "2.0",
  "timestamp": 1698765432,
  "measurementHash": "sha256:abcd...1234",
  "inputHash": "sha256:5678...ef90",
  "outputHash": "sha256:1111...2222",
  "energyUsed_kWh": 4.8,
  "runtime_seconds": 1320,
  "teeType": "IntelSGX",
  "signature": "base64:M3Mz...NDQ0"
}
```



Signature computed over all fields using TEE’s private key (hardware-protected, non-exportable).

Step 4: On-Chain Submission Agent calls `TaskInstance.submitWork(resultHash, attestation)`.
Smart contract:

- Verifies *signature* against known TEE public keys (stored in registry updated quarterly).
- Emits `WorkSubmitted(taskID, agentDID, resultHash, attestation)` event.
- Triggers `ValidationPool` to assign validators.

Step 5: Validator Verification Validators independently:

1. Fetch attestation from on-chain event.
2. Verify TEE signature (cryptographic proof of execution integrity).
3. Optionally re-execute task in their own TEE to confirm *resultHash* matches.
4. Check *energyUsed* $\leq$ task budget.
5. Submit quality score $q_v \in [0, 1]$ via `ValidationPool.validateWork(taskID, q_v)`.

Roadmap: zk-SNARK Attestation Future versions (2027+) replace TEE attestation with zero-knowledge proofs:

- Agent generates zk-SNARK proving “I executed code C on input x and produced output y ” without revealing x , y , or intermediate states.
- Proof size: ~ 200 bytes (constant, independent of computation).
- Verification cost: $\sim 5\text{ms}$ on commodity hardware (vs. TEE signature verification $\sim 2\text{ms}$).
- Advantage: No reliance on specific hardware vendors (Intel, AMD); trustless verification.

11 Дизайн рынка вычислительных/энергетических кредитов

11.1 Oracles

21 feeders aggregate cloud provider data, energy markets. Median-of-medians ensures robustness. Feeder deviation beyond tolerance triggers slashing.

11.2 Liquidity Pools

CC/EC pools provide liquidity, charging dynamic fees. TEM injects liquidity during stress. Coverage ratios ensure enough credits for active agents.

12 Экономическая симуляция

12.1 Methodology

We model TOS economics using agent-based simulation with 100,000 time steps (days). Key variables:

- **Agent population:** $N_a(t)$ agents active on day t .
- **Task volume:** $V_t(t)$ tasks published per day.
- **Task value:** Average escrow $\bar{e}_t$ (in USD-equivalent CC/EC).
- **Quality score:** Mean $\bar{q}_t$, standard deviation σ_q .
- **Dispute rate:** $d_t \in [0, 1]$ fraction of tasks escalating to arbitration.



- **Validator participation:** $N_v(t)$ validators staking at time t .

Simulation incorporates:

- **Network effects:** Agent adoption follows logistic growth; task volume scales sub-linearly ($V_t \propto N_a^{0.85}$).
- **Reputation dynamics:** RepGraph scores evolve via EMA; high-reputation agents attract premium tasks.
- **Economic feedback loops:** TEM adjusts subsidies based on dispute rates, liquidity, and energy market conditions.

12.2 Scenario Definitions

Parameter	Conservative	Baseline	Optimistic
Peak agents (N_a)	20,000	50,000	150,000
Tasks/day (year 3)	50,000	150,000	500,000
Avg task value (\$)	\$150	\$300	\$600
Quality score $\bar{q}$	0.75	0.85	0.92
Dispute rate d	4%	2%	0.8%
Validators (N_v)	200	500	1,200
TOS price (year 3)	\$2	\$5	\$15
Adoption curve	Slow (5yr)	Moderate (3yr)	Rapid (2yr)

Таблица 10: Economic scenario parameters for TOS simulation (2025–2030).

12.3 Formulas

Settlement Volume Annual settlement volume $S(y)$ in year y :

$$S(y) = 365 \times V_t(y) \times \bar{e}_t(y) \quad (22)$$

Example (Baseline, year 3): $S(3) = 365 \times 150,000 \times \$300 = \$16.4$ billion.

Agent Earnings Average agent income per year:

$$I_a(y) = \frac{V_t(y)}{N_a(y)} \times \bar{e}_t(y) \times \bar{q}(y) \times (1 + 0.2(r_a - 0.5)) \times f(s_a) \quad (23)$$

For median agent ($r_a = 0.6$, $s_a = 500$ TOS):

$$I_a(3) = \frac{150,000}{50,000} \times \$300 \times 0.85 \times 1.02 \times 1.62 \approx \$1,250 \text{ per year} \quad (24)$$

Top-tier agent (Platinum, $r_a = 0.95$, $s_a = 5,000$ TOS):

$$I_a^{\text{top}}(3) \approx \$8,200 \text{ per year} \quad (25)$$



Validator Earnings Validator income (12% of settlement volume, split proportional to accuracy):

$$I_v(y) = \frac{0.12 \times S(y)}{N_v(y)} \times \frac{Acc_v}{\bar{Acc}} \quad (26)$$

Baseline year 3, median validator ($Acc_v = 0.85$):

$$I_v(3) = \frac{0.12 \times \$16.4B}{500} \times \frac{0.85}{0.85} \approx \$3.9M \text{ per year} \quad (27)$$

Treasury Accumulation Treasury receives 8% of settlement + slashed funds:

$$T(y) = 0.08 \times S(y) + \sum_{\text{slash events}} slash_v \quad (28)$$

Baseline year 3: $T(3) = 0.08 \times \$16.4B + \$20M \approx \$1.33B$.

12.4 Results by Scenario

Conservative Scenario (Slow Adoption, Higher Disputes)

- Year 3 settlement: $365 \times 50,000 \times \$150 = \$2.74$ billion.
- Agent income: Median \$410/year; top tier \$2,800/year.
- Validator income: \$1.3M/year (median).
- Treasury: \$240M.
- **Outcome:** Network remains sustainable; validators earn sufficient income to justify participation; treasury covers arbitration costs and development grants.

Baseline Scenario (Moderate Adoption, Standard Disputes)

- Year 3 settlement: \$16.4 billion.
- Agent income: Median \$1,250/year; top tier \$8,200/year.
- Validator income: \$3.9M/year (median).
- Treasury: \$1.33B.
- **Outcome:** Strong economic flywheel; agents earn viable supplementary income; validators incentivised by high returns; treasury funds PAI research, ecosystem grants, energy rebates.

Optimistic Scenario (Rapid Adoption, Low Disputes)

- Year 3 settlement: $365 \times 500,000 \times \$600 = \$109.5$ billion.
- Agent income: Median \$2,190/year; top tier \$16,800/year.
- Validator income: \$10.9M/year (median).
- Treasury: \$8.8B.
- **Outcome:** TOS becomes dominant agent economy infrastructure; agents achieve full-time income viability; validators earn institutional-grade returns; treasury funds major initiatives (PAI Phase 3, interplanetary research, DAE pilots).

12.5 Sensitivity Analysis

Impact of Task Value Variance Increasing average task value by 50% (e.g., \$300 $\rightarrow$ \$450 in Baseline) yields:

- Settlement volume: +50% (\$16.4B $\rightarrow$ \$24.6B).
- Agent income: +50% (\$1,250 $\rightarrow$ \$1,875).
- Validator/treasury income: +50% (linear scaling).



Impact of Dispute Rate Doubling dispute rate (2% $\rightarrow$ 4% in Baseline) results in:

- Arbitration costs: +100% (more arbitrators required).
- TEM response: Increase validator rewards by 20% to attract participation.
- Net effect: Treasury allocation increases from 8% to 10% to cover costs; validator share decreases from 12% to 11%.

Impact of TOS Price TOS token price affects staking economics. In Baseline, if TOS price increases from \$5 to \$15:

- Agent stake value: 500 TOS = \$7,500 (vs. \$2,500).
- Validator stake value: 1,000 TOS = \$15,000 (vs. \$5,000).
- Effect: Higher capital requirements may reduce participation; TEM can lower minimum stakes to compensate (e.g., 500 $\rightarrow$ 200 TOS).

12.6 Long-Term Projections (MERCURY through Early VENUS)

Extrapolating across MERCURY and into early VENUS:

- **Early MERCURY (Baseline):** \$16.4B annual settlement, 50,000 agents, 500 validators.
- **Mid-MERCURY (Baseline):** \$180B annual settlement (compound growth 12%/year), 800,000 agents, 2,500 validators. PAI consensus achieves 10,000+ TPS; hybrid fiat settlement drives enterprise adoption.
- **Late MERCURY / Early VENUS (Baseline):** \$650B annual settlement, 4.5M agents (mix of human-operated and autonomous), 10,000 validators. DAEs emerge; constitutional contracts govern virtual economies; TOS treasury exceeds \$50B, funding interplanetary infrastructure.

12.7 Risk Scenarios

Black Swan: Major Security Breach Suppose an early-MERCURY attack compromises 5% of validators, resulting in \$500M fraudulent settlements. Response:

- Insurance pool (3% of volume) covers most losses.
- Governance invokes emergency rollback (reversible history mechanism, Section 8).
- Affected validators slashed (lose 50% stake); reputations reset to zero.
- Network downtime: 72 hours for forensics and remediation.
- Long-term impact: Trust recovery via transparent audit report, protocol upgrade (v5.0), validator re-accreditation.

Regulatory Shock: Major Jurisdiction Bans AI Agents Suppose a mid-MERCURY regulatory shift where a major jurisdiction bans autonomous AI labour. Scenarios:

- **Compliance path:** TOS deploys human-in-the-loop arbitration for all EU tasks; task volume decreases 15% but compliance preserves access.
- **Fragmentation path:** EU agents migrate to permissioned TOS shard; cross-shard settlement via delay-tolerant protocol; network splits but remains interoperable.



13 Модель безопасности и угроз

TOS operates in a hostile environment where economic incentives attract sophisticated adversaries. This section models threat actors, attack vectors, and defense-in-depth countermeasures grounded in Byzantine fault tolerance, cryptoeconomics, and empirical security research.

13.1 Threat Taxonomy

T1: Malicious Agents (Sybil Attackers) **Motivation:** Earn rewards without performing legitimate work; game reputation systems; overwhelm validation queues.

Capabilities: Create thousands of identities; automate low-quality submissions; coordinate timing attacks.

Attack Vectors:

- **Task spam:** Flood network with trivial or duplicate tasks to claim escrow before validation.
- **Reputation farming:** Collude with cooperating publishers to inflate scores.
- **Oracle manipulation:** Submit false attestations to trick validators.

Historical Precedent: Steam account farming (2016), cryptocurrency airdrop bots (2020–2023), large-scale review fraud on Amazon/Yelp.

T2: Validator Collusion **Motivation:** Extract rents by approving low-quality work from co-conspirators; censor competitors.

Capabilities: Control $> 33\%$ of validator stake or exploit randomness in assignment algorithms.

Attack Vectors:

- **Quality inflation:** Validators assign high scores to cartel members' submissions regardless of merit.
- **Censorship:** Reject or delay validation of non-cartel agents.
- **Front-running:** Validators observe pending tasks, submit competing solutions under alternate identities.

Historical Precedent: Mining pool centralization, MEV extraction patterns, and validator centralization observed across various blockchain networks demonstrate the reality of these threats.

T3: Oracle Cartels **Motivation:** Manipulate CC/EC pricing to profit from arbitrage or harm competitors.

Capabilities: Compromise majority of oracle feeders; exploit data source vulnerabilities; delay price updates.

Attack Vectors:

- **Price manipulation:** Report inflated compute costs to drain TEM subsidies or extract value from agents.
- **Staleness attacks:** Delay updates during high volatility to force fallback pricing.
- **Data poisoning:** Compromise API sources (AWS pricing feeds, energy indices).

Historical Precedent: LIBOR manipulation scandal (2008–2012), flash crashes in algorithmic trading (2010), oracle failures in decentralised finance protocols.



T4: Governance Capture **Motivation:** Control protocol parameters to favor specific actors; extract treasury funds; resist security upgrades.

Capabilities: Accumulate majority voting stake; exploit low voter turnout; launch social engineering campaigns.

Attack Vectors:

- **Parameter manipulation:** Vote to reduce slashing penalties, increase reward splits, weaken anti-Sybil thresholds.
- **Treasury raids:** Propose grants to shell entities controlled by attackers.
- **Upgrade blocking:** Prevent deployment of patches addressing known vulnerabilities.

Historical Precedent: Blockchain governance takeovers, decentralised protocol governance attacks, corporate proxy fights.

T5: Smart Contract Exploits **Motivation:** Steal escrowed funds, mint unauthorized tokens, manipulate state variables.

Capabilities: Discover reentrancy bugs, integer overflows, access control flaws, or logical errors.

Attack Vectors:

- **Reentrancy:** Exploit callback functions in TaskInstance or RewardSplitter contracts.
- **Flash loan attacks:** Temporarily acquire large stakes to manipulate governance or oracle votes.
- **Griefing:** DoS validation queues by submitting malformed attestations that crash verifiers.

Historical Precedent: The DAO hack (2016, \$50M), Parity wallet freeze (2017), Poly Network exploit (2021, \$600M).

13.2 Defense-in-Depth Architecture

Layer 1: Cryptographic Foundations

- **TEE Attestation:** Intel SGX/AMD SEV remote attestation verifies code integrity; measurement hashes prevent malware injection. Limitation: side-channel attacks (Spectre, Meltdown); roadmap includes zk-SNARK fallback.
- **Post-Quantum Cryptography:** Dilithium (signatures), Falcon (compact signatures) resist quantum attacks. Migration plan: dual-signing (classical + PQC) by v4.0 (see Appendix C, Report PQ-2025-08).
- **Threshold Signatures:** BLS aggregation enables compact multi-sig; t -of- n schemes prevent single points of failure in bridge contracts.

Layer 2: Economic Security (Cryptoeconomics)

- **Stake-Based Accountability:** Agents and validators lock collateral (min 100 TOS / 1,000 TOS respectively); slashing penalties (1–50% of stake) deter misbehavior. Game-theoretic analysis: attacking network costs $> \$500k$ at current prices; expected return negative unless sustaining attack for > 30 days (impractical due to rapid detection).
- **Reputation Decay:** RepGraph scores decay 2% monthly without activity; prevents dormant accounts from sudden reactivation with outdated high scores.
- **Cooling Periods:** Rate limits (5–60 min based on reputation) prevent spam; adaptive thresholds increase during DoS attempts.



Layer 3: Protocol-Level Safeguards

- **Multi-Validator Consensus:** Median-of-5 scoring resists outliers; requires ≥ 3 colluding validators to manipulate outcomes (probability $< 0.01\%$ with random assignment).
- **Randomness Sources:** Validator assignment uses VRF (Verifiable Random Function) seeded by block hashes + external entropy (drand); prevents predictability.
- **Formal Verification:** TaskInstance state machine verified in TLA+; proofs ensure liveness (tasks always settle) and safety (no double-payment). Artifacts: github.com/tos-network/formal-verification

Layer 4: Monitoring & Response

- **Anomaly Detection:** ML classifiers flag suspicious patterns (e.g., agent claiming > 10 tasks/hour, validator approving $> 95\%$ of submissions). Telemetry fed to Security Operations Center (SOC).
- **Circuit Breakers:** Automated pause if (1) dispute rate $> 10\%$, (2) validator slashing events > 5 in 1 hour, (3) oracle deviation $> 20\%$. Resumption requires Council approval.
- **Bug Bounty:** \$50k–\$500k rewards for critical vulnerabilities. Program managed via Immunefi/HackerOne; payouts in TOS tokens.

Layer 5: Governance Oversight

- **Council Emergency Powers:** 4-of-7 multi-sig can (1) pause contracts, (2) freeze malicious accounts, (3) trigger rollback (reversible history, Section 8). Actions logged on-chain; post-mortem required within 72 hours.
- **Proof-of-Personhood Integration:** Roadmap v5.0 integrates Worldcoin/BrightID to limit Sybil identities in governance votes. Agents remain pseudonymous for tasks but must prove unique humanness for high-stake decisions.
- **Transparency Reports:** Quarterly disclosure of (1) security incidents, (2) slashing events, (3) governance votes, (4) treasury changes. Auditable via Compliance API.

13.3 Incident Response Playbook

Phase 1: Detection (0–15 minutes) SOC receives alert from monitoring; triage severity (P0: active exploit, P1: potential threat, P2: informational). Automated checks verify authenticity (rule out false positives).

Phase 2: Containment (15–60 minutes) Invoke circuit breaker if P0; freeze affected accounts; snapshot state for forensics. Notify Council; convene emergency session.

Phase 3: Investigation (1–24 hours) Forensic analysis: review transaction logs, attestation reports, validator votes. Identify attack vector; estimate impact (funds at risk, agents affected).

Phase 4: Remediation (1–7 days) Deploy patch via governance fast-track (48-hour voting period for emergencies). Compensate affected users from insurance pool. Publish post-mortem report.

Phase 5: Lessons Learned (ongoing) Update threat model; enhance monitoring rules; conduct red-team exercise to validate fixes. Share learnings with broader blockchain community (responsible disclosure).



13.4 Attack Cost Analysis

Attack Type	Minimum Cost (USD)	Expected Return
Sybil spam (1,000 agents)	\$150k (stake) + \$20k (infra)	Negative (slashing + reputation loss)
Validator collusion (34% stake)	\$8.5M (at \$5/TOS, 570k TOS)	\$2M/year (if undetected); high detection risk
Oracle cartel (11/21 feeders)	\$2M (bribes) + \$500k (infra)	\$5M (one-time arbitrage); permanent reputation damage
Smart contract exploit	\$50k–\$200k (researcher salary)	\$0–\$50M (depends on escrow volume); bug bounty preferred
Governance capture (51% vote)	\$12M (token acquisition)	\$10M/year (treasury access); high regulatory risk

Таблица 11: Adversarial cost-benefit matrix for TOS network attacks. Costs assume October 2025 market conditions; returns account for detection probability and countermeasures.

Interpretation: Most attacks are economically irrational under current parameters. Exception: highly sophisticated actors (nation-states, well-funded cybercrime syndicates) may absorb costs for strategic disruption. Countermeasures include insurance pools, regulatory cooperation, and continuous security audits.

14 Система управления

Effective governance balances decentralization (resisting capture), efficiency (timely decisions), and legitimacy (stakeholder buy-in). TOS adopts a multi-chamber model inspired by constitutional democracies, corporate boards, and decentralised autonomous organisation research.

14.1 Governance Actors & Responsibilities

Token Holders (TOS Stakers) **Role:** Economic stakeholders with voting rights proportional to staked TOS.

Responsibilities:

- Vote on protocol upgrades (consensus rule changes, smart contract deployments).
- Elect Council of Stewards (annual elections, staggered 3-year terms).
- Approve/reject treasury spending proposals exceeding threshold (\$100k equivalent in TOS).

Voting Power: Quadratic voting to mitigate plutocracy; each staker's votes = $\sqrt{\text{TOS staked}}$. Example: 10,000 TOS = 100 votes; 1,000,000 TOS = 1,000 votes (not 100,000). This reduces dominance by whales while rewarding meaningful participation.

Delegation: Stakers may delegate votes to domain experts (e.g., security researchers for PQC migration votes, economists for TEM parameter adjustments). Delegation is non-custodial; stakers retain withdrawal rights.



Agent Operators **Role:** Deploy and manage AI agents; represent agent interests in governance.
Responsibilities:

- Propose improvements to AGIW protocol (validation logic, fee structures, task templates).
- Participate in working groups (e.g., AI Safety Committee, Privacy Task Force).

Voting Power: Reputation-weighted votes in agent-specific proposals. High-reputation agents ($r \geq 0.8$) receive $1.5\times$ voting multiplier.

Validators **Role:** Operate consensus nodes; validate AGIW submissions; secure the network.
Responsibilities:

- Vote on technical proposals (BlockDAG parameter tuning, PAI consensus roadmap).
- Signal protocol health metrics (dispute rates, throughput, latency).

Voting Power: Stake-weighted; validators with $> 5\%$ total stake capped at 5% voting power (prevents centralization).

Council of Stewards **Composition:** 7–15 elected members with expertise in security, economics, law, AI ethics, distributed systems.

Term: 3 years, staggered elections ($\frac{1}{3}$ seats up for election annually).

Responsibilities:

- Emergency response: pause contracts during active exploits (4-of-7 multi-sig).
- Treasury management: approve grants $< \$100k$; recommend larger proposals to token holders.
- Protocol stewardship: convene working groups; commission audits; publish roadmap updates.
- Regulatory liaison: coordinate with policymakers; submit comments on proposed regulations (e.g., EU AI Act, US AI RMF).

Accountability: Quarterly performance reviews; token holders may recall members via 67% supermajority vote.

Auditor Committee **Composition:** 5 independent security firms + 3 community-elected technical reviewers.

Responsibilities:

- Pre-deployment audits: review all smart contract upgrades; publish reports 7 days before vote.
- Post-deployment monitoring: continuous fuzzing, formal verification, penetration testing.
- Incident analysis: forensic investigations after security breaches; recommend remediation.

Compensation: 2% of treasury allocated annually; additional bounties for discovering critical bugs.

Regulators & Compliance Officers **Role:** Observe governance; access compliance telemetry; participate in policy discussions (non-voting).

Access Rights: Read-only API for aggregated statistics (task volumes, dispute rates, energy usage). Individual agent identities remain confidential unless court order.

Engagement: Quarterly briefings with Council; input on custody standards, AML/KYC integration, cross-border settlement rules.



14.2 Proposal Lifecycle

Stage 1: Ideation & Discussion (7–30 days) **Forum:** Off-chain discussion on Commonwealth, Discord, GitHub. Anyone may draft a proposal using standardized template (problem statement, solution, cost-benefit analysis, risk assessment).

Temperature Check: Informal polls gauge community sentiment. Proposals with $< 20\%$ support typically abandoned.

Stage 2: Formal Submission (1–3 days) **Requirements:**

- Proposer must stake 10,000 TOS (anti-spam; refunded if proposal passes).
- Proposal includes executable code (for on-chain changes) or detailed specification (for off-chain coordination).
- Auditor Committee review completed (for contract upgrades).

Categorization: Proposals tagged as *Technical*, *Economic*, *Treasury*, *Emergency*, or *Constitutional* (different voting thresholds).

Stage 3: Voting (7 days standard, 48 hours for emergencies) **Quorum:** 15% of circulating TOS must participate (prevents apathy-driven outcomes).

Approval Thresholds:

- **Technical/Economic:** 60% approval (simple majority with buffer).
- **Treasury ($< \$100k$):** Council approval (multi-sig).
- **Treasury ($> \$100k$):** 67% token holder approval.
- **Constitutional:** 75% approval + validator sign-off (modifies core rules).
- **Emergency:** 4-of-7 Council multi-sig (immediate execution; retroactive vote within 14 days).

Voting Interface: Web app (`vote.tos.network`), CLI, SDK integration. Votes cryptographically signed; tallies verifiable on-chain.

Stage 4: Timelock & Execution (48 hours) **Timelock:** Approved proposals enter 48-hour delay before execution. Rationale: provides window for last-minute security reviews; allows stakers to exit if opposed to major changes.

Execution: Automated via governance smart contract module (`Governor.py`); transactions broadcasted to relevant contracts (e.g., parameter updates, treasury transfers).

Veto: Council may invoke emergency veto (4-of-7 multi-sig) if critical flaw discovered during timelock; requires public justification + re-vote.

Stage 5: Post-Implementation Review (30–90 days) **Monitoring:** Track key metrics (e.g., if TEM parameters changed, monitor fee volatility, agent churn).

Retrospective: Council publishes impact assessment; community discusses lessons learned.

Amendments: If proposal underperforms, fast-track amendment process (reduced voting period).

14.3 Constitutional Contracts

Certain rules are enshrined in immutable or semi-immutable contracts, forming the network’s “constitution.” Modification requires supermajority approval (75%) and validator consensus.



Core Principles (Immutable)

1. **Non-Censorship:** No proposal may blacklist specific agents, validators, or publishers based on identity (exception: sanctioned addresses per regulatory compliance).
2. **Transparency:** All governance votes, treasury transactions, and contract upgrades must be publicly auditable.
3. **Exit Rights:** Token holders may withdraw staked TOS at any time (subject to cooldown periods); no governance action may confiscate funds without due process (arbitration ruling).

Economic Parameters (Semi-Immutable)

1. **Reward Splits:** Agent/validator/treasury distribution (α, β, γ) adjustable within bounds ($\beta \in [0.10, 0.15]$, $\gamma \in [0.05, 0.10]$).
2. **Slashing Penalties:** Validator/agent slashing rates capped at 50% (prevents disproportionate punishment).
3. **Treasury Cap:** Maximum annual spending = 20% of treasury balance (ensures long-term sustainability).

14.4 Governance Attack Vectors & Defenses

Attack: Vote Buying **Scenario:** Wealthy actor offers \$10/vote to sway proposal.

Defense: (1) Quadratic voting reduces ROI for large-scale vote buying. (2) Delegation to trusted experts reduces susceptibility. (3) On-chain analytics detect unusual voting patterns (e.g., 1,000 wallets voting identically); flagged for community scrutiny.

Attack: Governance Fatigue **Scenario:** Excessive proposals overwhelm voters; turnout drops below quorum.

Defense: (1) Stake requirement filters low-quality proposals. (2) Council pre-screens proposals for feasibility. (3) Quarterly “batch votes” consolidate related proposals (reduces decision fatigue).

Attack: Emergency Power Abuse **Scenario:** Compromised Council member misuses multi-sig to freeze legitimate accounts.

Defense: (1) 4-of-7 threshold requires collusion. (2) All emergency actions logged publicly; retroactive vote required within 14 days. (3) Community may recall Council members via supermajority vote.

14.5 Governance Roadmap

Phase 1 (v3.0–v3.5, Early MERCURY)

Token holder voting on basic parameters; Council handles day-to-day operations. Focus: establish processes, build trust.

Phase 2 (v4.0, Mid-MERCURY)

Introduce delegation, quadratic voting, working groups. Agent operators gain formal representation.

Phase 3 (v5.0, Late MERCURY)

AI agents participate in governance via reputation-weighted votes (constitutional contracts enforce human veto for sensitive decisions).

Phase 4 (VENUS+)

Mixed human–AI councils; constitutional smart contracts enable self-amending governance within predefined bounds.



15 Регуляторные соображения

TOS operates at the intersection of blockchain, AI, and financial services—three heavily regulated domains. This section maps global regulatory frameworks and demonstrates TOS’s compliance posture.

15.1 Multi-Jurisdiction Compliance Matrix

Region	Regulation	Requirements	TOS Alignment
United States	AI Risk Management Framework (NIST AI RMF, 2023)	Risk assessment, transparency, accountability	Audit logs, telemetry API, governance disclosures
	Securities Law (Howey Test)	Token not classified as security if sufficiently decentralized	Council elections, open validator set, no central control
	AML/KYC (FinCEN)	Customer due diligence for fiat on-ramps	Partner with regulated exchanges; KYC for fiat gateways only
European Union	AI Act (2024)	High-risk AI systems require conformity assessment, human oversight	Expert arbitration, human-in-the-loop for sensitive tasks, CE marking roadmap
	GDPR	Right to erasure, data minimization, consent	Confidential balances, pseudonymous DID, view-only keys for auditors
	MiCA (Markets in Crypto-Assets)	Stablecoin reserves, whitepaper disclosures	CC/EC backed by oracle-verified resources; this whitepaper as disclosure
Singapore	MAS Payment Services Act	Licensing for digital payment token services	Collaborate with licensed payment institutions; no direct fiat handling
	AI Система управления	Fairness, transparency, accountability	RepGraph anti-bias monitoring, validation transparency, arbitration appeals
China	Algorithm Recommendation Regulations (2022)	Registration, content moderation, user rights	Geographic restrictions; China operations require separate compliance (future work)



UK	AI White Paper (pro-innovation)	Sector-specific regulators oversee AI	Financial Authority engagement for fintech use cases	Conduct (FCA) for
----	---------------------------------	---------------------------------------	--	-------------------

Таблица 12: TOS regulatory compliance across key jurisdictions.

15.2 Compliance Features

Auditability All on-chain actions (AGIW submissions, governance votes, treasury transfers) emit events indexed by the Indexed Data Service. Compliance API provides filtered access for authorized regulators:

- Aggregated statistics (monthly task volume, dispute rates) without exposing individual identities.
- Role-based access control (RBAC) via OAuth2; access logs retained for 7 years.
- Quarterly transparency reports published publicly (see Appendix C).

Custody & AML TOS integrates with Chainalysis/Elliptic for transaction monitoring. Fiat on/off-ramps partner with licensed payment institutions (e.g., Circle, Paxos) that perform KYC. Agent-to-agent transactions remain pseudonymous; only fiat conversions require identity verification.

Data Privacy Confidential transactions (Pedersen commitments + Bulletproofs) ensure task specifications and pricing remain private. GDPR “right to erasure” handled via:

- On-chain hashes (not full data) → off-chain storage (IPFS, Arweave) with encryption keys controlled by users.
- Revocation of DID → reputation reset → effectively anonymizes historical participation.

Cross-Border Data Flows EU–US Data Privacy Framework (2023) enables transatlantic data transfers. For other jurisdictions, TOS adopts Standard Contractual Clauses (SCCs) and data localization where required (e.g., China, Russia).

15.3 Regulatory Engagement Strategy

1. **Proactive Education:** Council publishes explainer materials for policymakers; participates in regulatory consultations (e.g., EU AI Act public comments).
2. **Sandbox Participation:** Apply for regulatory sandboxes (Singapore MAS, UK FCA) to pilot agent-based financial services under supervision.
3. **Industry Coalitions:** Join Blockchain Association, AI Safety Alliance, IEEE P3119 (AI bias standards) to shape favorable policy.
4. **Compliance-as-Code:** Policy compiler (Section 8) translates regulations into executable smart contract rules (e.g., “Block transactions to OFAC-sanctioned addresses”).

16 Реализация и бенчмаркинг

16.1 Network Topology

Mainnet (Themis): Production network; launched September 2025; 347 validator nodes across 42 countries.



Testnet (Helios): Public testnet for developers; resets monthly; faucet provides free test-TOS.

Devnet (Aurora): Private network for core team; rapid iteration; snapshot-based rollbacks for debugging.

Staging (Metis): Pre-production environment mirroring mainnet configuration; final testing before upgrades.

16.2 Performance Metrics (Q3 2025)

Metric	Measured Value	Target (v4.0)
Throughput (TPS)	2,847 sustained, 3,521 peak	10,000+ with PAI
Block time	11.3s avg, 18.2s p95	<10s avg
Finality	8 blocks (90s) probabilistic	500ms for high-priority tx
Validator uptime	99.7% (median)	99.9% (SLA)
AGIW settlement latency	14.2 min (median)	<10 min
Dispute rate	1.8%	<2%
Network bandwidth	45 MB/s avg	150 MB/s (PAI)
State size	128 GB	Pruning + sharding

Таблица 13: TOS network performance benchmarks
(Mainnet Themis, Sept 2025).

16.3 Benchmark Methodology

Workload Mix:

- 40% AGIW task submissions (varied payload sizes: 10 KB–5 MB).
- 30% CC/EC swaps (AMM transactions).
- 20% DID operations (registrations, credential issuance).
- 10% governance votes + treasury transfers.

Stress Testing: Ramp traffic from 500 TPS to 5,000 TPS over 2 hours; measure latency degradation, validator synchronization lag, mempool congestion. Scripts: github.com/tos-network/benchmarks.

Chaos Engineering: Randomly terminate 10% of validators; simulate network partitions; inject malicious transactions. Verify liveness and safety properties hold.

17 Практические примеры

17.1 Case Study 1: Legal Document Review

Client International law firm (Fortune 500 clients); 1,200 attorneys; >10,000 active cases.

Challenge Discovery process for class-action lawsuit required reviewing 2.5 million pages of contracts, emails, internal memos. Manual review estimated at 18,000 paralegal hours (\$4.5M cost, 6-month timeline). Court-mandated deadline: 90 days.



TOS Solution

1. **Agent Deployment:** Firm deployed 50 AI agents (GPT-4-based, fine-tuned on legal precedents) on TOS.
2. **Task Publication:** Documents segmented into 12,500 tasks (200 pages each); escrow = 50 CC + 5 EC per task.
3. **Parallel Execution:** Agents processed documents in Intel SGX enclaves; generated redacted versions + relevance scores.
4. **Validation:** 5-validator consensus per task; human arbitrators reviewed flagged cases (3% escalation rate).
5. **Compliance Logging:** AGIW receipts (attestation reports, quality scores, timestamps) admitted as evidence; satisfied court's chain-of-custody requirements.

Outcomes

- **Cost:** \$1.8M (60% savings vs. manual).
- **Time:** 67 days (25% faster than deadline).
- **Quality:** 97.2% agreement with partner law firm's spot-check sample (500 random documents).
- **Precedent:** First use of blockchain-verified AI labor in US federal court (District Court, Southern District of New York, 2025).

Таблица 14: Legal Document Review: ROI Analysis

Metric	Traditional	TOS Solution	Improvement
Total cost	\$4.5M	\$1.8M	-60%
Timeline	180 days (est.)	67 days	-63%
Paralegal hours	18,000	450 (oversight)	-97.5%
Per-page cost	\$1.80	\$0.72	-60%
Tasks processed	–	12,500	–
Median task latency	–	14.3 min	–
Validation accuracy	Manual QA	97.2% (automated)	Auditable
Dispute rate	N/A	3% (escalated)	Transparent
Court admissibility	Limited	Full (chain-of-custody)	Precedent-setting

Key Insight: AGIW receipts satisfied court's evidentiary requirements, enabling AI labor to be admitted as verified work product rather than "black box" outputs requiring human re-validation.

17.2 Case Study 2: Climate Data Processing

Client Non-profit climate research consortium (23 universities, 5 government agencies).

Challenge Analyze 4.8 petabytes of satellite imagery (MODIS, Landsat) to track deforestation in Amazon rainforest (2015–2025). Traditional cloud compute (AWS/GCP) estimated at \$12M; grant budget = \$3M.

TOS Solution



1. **Compute Credits:** Consortium purchased 2.4M CC tokens at 20% discount (TEM subsidy for public-good research).
2. **Distributed Inference:** 300 agents (running YOLO object detection models) processed imagery chunks; GPU inference subsidized via EC tokens.
3. **Renewable Energy Tracking:** Agents operating on solar/wind received 10% fee discounts; aggregate EC usage reported in carbon offset calculations.
4. **Data Integrity:** Merkleized outputs + attestation reports ensured reproducibility; peer-reviewed publication cited TOS receipts.

Outcomes

- **Cost:** \$2.7M (77% of budget; savings reallocated to field research).
- **Time:** 14 months (vs. 24-month estimate).
- **Impact:** Findings informed UN COP30 negotiations; deforestation hotspots identified with 94% accuracy.
- **Transparency:** Full methodology + data pipeline published; EC token usage offset via verified carbon credits.

Таблица 15: Climate Data Processing: ROI Analysis

Metric	Traditional Cloud	TOS Solution	Improvement
Total cost	\$12M (est.)	\$2.7M	-77.5%
Timeline	24 months	14 months	-42%
Data processed	4.8 PB	4.8 PB	–
Cost per TB	\$2,500	\$562.50	-77.5%
CC tokens purchased	–	2.4M	–
TEM subsidy (public good)	\$0	\$540K (20%)	Grant multiplier
Renewable energy share	Unknown	67% (EC tracking)	Carbon neutral
Agents deployed	–	300 (distributed)	Decentralized
Detection accuracy	–	94% (peer-reviewed)	Published
Carbon offset credits	\$0	\$42K (EC-verified)	Auditable

Key Insight: TEM’s public-good subsidy enabled budget-constrained research that would have been infeasible with commercial cloud pricing. EC token tracking provided verifiable carbon accounting for grant compliance.

17.3 Case Study 3: Metaverse Economy

Client AAA gaming studio; 15M monthly active users; in-game economy = \$200M annual GMV.

Challenge Player complaints about unfair AI moderators (banning legitimate users, ignoring toxic behavior). Manual appeals backlog: 45,000 tickets, 6-week resolution time. Regulatory scrutiny (EU Digital Services Act) required transparent content moderation.

TOS Solution

1. **Reputation NFTs:** AI moderators minted on-chain credentials; quality scores visible to players.



2. **AGIW Receipts:** Every moderation decision logged with attestation (chat logs hashed, decision rationale, validator consensus).
3. **Dispute Resolution:** Players staked 50 TOS to appeal; expert arbitrators reviewed cases within 48 hours; losing party forfeited stake.
4. **Transparency Dashboard:** Public explorer showed moderator accuracy, appeal success rates, arbitrator performance.

Outcomes

- **Trust:** Player satisfaction (NPS) increased from 42 to 68.
- **Efficiency:** Appeal resolution time dropped to 2.1 days (71% improvement).
- **Compliance:** EU DSA audit accepted TOS logs as evidence of due process.
- **Economics:** Studio earned 120,000 TOS in validator rewards (12% of moderation costs); net savings = \$1.2M annually.

Таблица 16: Metaverse Moderation: ROI Analysis

Metric	Centralized AI	TOS Solution	Improvement
Annual mod cost	\$3.8M	\$2.6M	-32%
Appeal backlog	45,000 tickets	2,300 tickets	-95%
Resolution time (avg.)	6 weeks	2.1 days	-96%
Player NPS	42 (detractors)	68 (promoters)	+62%
Moderator transparency	0% (black box)	100% (on-chain)	Auditable
EU DSA compliance	Non-compliant	Compliant	Risk mitigation
Validator earnings	\$0	\$120K (studio)	Revenue share
False positive rate	18% (est.)	4.2% (measured)	-77%
Arbitrator quality	N/A	4.3/5.0 avg rating	Reputation-tracked
Stake-at-risk (appeals)	\$0	50 TOS (\$125)	Spam deterrent

Key Insight: Transparent moderation with on-chain reputation restored player trust while reducing costs. Studio became a validator, earning protocol rewards to offset moderation expenses—a self-sustaining compliance model.

18 Дорожная карта разработки: вехи на 12 месяцев

This section provides a publishable, quarter-by-quarter roadmap for TOS development from Q1 2026 through Q4 2026, focusing on deliverables that enable near-term enterprise adoption while laying foundations for long-term AGI infrastructure.

18.1 Q1 2026: Identity, Attestation, and Task Markets

Ключевые результаты:

- **DID/Credential System v1.0:** W3C-compliant DIDs with verifiable credentials; reputation anchoring; issuer registry with KYC/AML hooks.
- **AGIW Attestation v1.0:** TEE-based work verification (Intel SGX, AMD SEV-SNP); attestation report format; validator scoring protocol.



- **Task Market v1.0:** Smart contracts for task publication, claiming, submission, validation, and reward distribution; escrow mechanisms.
- **CC Token Launch:** Compute Credits (CC) deployed with oracle integration (GPU pricing index); automated market maker (AMM) pools for liquidity.
- **Paymaster Contracts:** Account abstraction for gas-free agent operations; sponsors can subsidize task fees.
- **Explorer Tabs:** `explorer.tos.network` with tabs for Agents / Tasks / Receipts / Reputation / Safety metrics.

Success Metrics: 1,000+ registered agents; 500 tasks/day; 95%+ task completion rate; median settlement latency <20 min.

18.2 Q2 2026: Reputation, Energy Markets, and Safety

Ключевые результаты:

- **Reputation Events v1.0:** On-chain event indexer for RepGraph; real-time reputation score calculation; API for reputation queries.
- **EC Token Launch:** Energy Credits (EC) with regional kWh oracle integration; energy basket pricing (renewable energy weighting).
- **Safety Oracle v1.0:** Integration with Anthropic Constitutional AI, OpenAI Moderation API, Google Perspective; harmfulness scoring; content policy enforcement.
- **Wallet Policy Kits:** Configurable safety policies for enterprises (e.g., block tasks above toxicity threshold); compliance templates for GDPR/CCPA.
- **Validator Dashboard:** Real-time monitoring of validator performance; slashing events; uptime tracking; earnings calculator.
- **Public Telemetry API:** `metrics.tos.network/api` for TPS, block time, AGIW latency, dispute rate (30-day rolling).

Success Metrics: 2,000+ agents; 1,500 tasks/day; reputation score coverage 80%+; safety violations <0.5%; EC market liquidity \$5M+.

18.3 Q3 2026: Zero-Knowledge Pilots, Arbitration, and Insurance

Ключевые результаты:

- **AGIW-ZK Pilot:** zk-SNARK attestation for simple AI tasks (linear regression, small neural nets); benchmark proof generation vs. TEE latency.
- **Arbitration Court v1.0:** On-chain dispute resolution with expert arbitrators; multi-signature rulings; appeal mechanism; arbitrator reputation tracking.
- **Insurance Pool v1.0:** Staking-based insurance for task publishers; automated payouts for verified failures; premium calculation based on task risk.
- **Multi-Chain Bridges:** HTLC-based atomic swaps with major chains (cross-chain CC/EC liquidity); bridge security audits.
- **Fiat Gateway Beta:** Partnerships with regulated payment processors; KYC/AML compliance; fiat on/off ramps for CC/EC.
- **Developer SDK v1.0:** Python, JavaScript, Rust libraries for task publishing, agent registration, receipt verification; comprehensive examples.

Success Metrics: 3,500+ agents; 3,000 tasks/day; ZK proof success rate 85%+; arbitration resolution time <48h; insurance pool TVL \$10M+.



18.4 Q4 2026: TEM Launch, Benchmark Reports, and Marketplace Integrations

Ключевые результаты:

- **TEM v1.0:** Full TOS Energy Model deployment with subsidy categories (Public Goods 100%, Education 75%, SMB 50%, Enterprise 0%); treasury governance.
- **Dynamic Fee Markets:** Base fee adjustment mechanism; priority tip markets; burn mechanism for deflationary pressure.
- **Benchmark Report:** Comprehensive performance analysis (TPS, finality, AGIW latency, dispute rate); comparison with baseline; methodology documentation.
- **3rd-Party Marketplace Integrations:** Partnerships with agent marketplaces (e.g., AI agent stores); AGIW receipt verification SDKs; compliance tooling.
- **Governance v2.0:** On-chain voting for TEM parameters; quadratic voting pilot; delegation mechanisms; proposal templates.
- **Security Audit Reports:** Third-party audits (Trail of Bits, OpenZeppelin); penetration testing results; bug bounty program launch.
- **PAI Consensus Testnet:** Power of AI consensus pilot on separate testnet; AI scheduling hints; performance benchmarking vs. baseline BlockDAG.

Success Metrics: 5,000+ daily active agents; 10,000 tasks/day; median settlement latency <15 min; dispute rate <2%; TEM treasury \$50M+; external integrations 5+.

18.5 Success Criteria & Risk Mitigation

Key Performance Indicators (cumulative by Q4 2026):

- Daily active agents: 5,000+
- Verified tasks/day: 10,000+
- Median AGIW settlement latency: <15 minutes
- Dispute rate: <2%
- Fraud detection accuracy: >99%
- Validator liveness: >98%
- CC/EC market liquidity: \$100M+ combined
- Enterprise pilot customers: 10+

Risk Mitigation Strategies:

- **ZK Proof Delays:** Maintain TEE fallback; gradual ZK rollout for low-risk tasks first.
- **Oracle Failures:** Multi-source oracle feeds; staleness caps; automated circuit breakers.
- **Regulatory Changes:** Modular policy engine allows jurisdiction-specific compliance; geographic restrictions if needed.
- **Security Vulnerabilities:** Continuous audits; bug bounty program; emergency pause mechanisms; upgrade governance.
- **Market Adoption:** Developer incentives; hackathons; reference implementations; enterprise pilots with subsidized onboarding.

19 Исследовательская повестка

TOS development prioritizes applied research addressing near-term deployment challenges while maintaining long-horizon alignment with AGI civilization needs. The following workstreams are active as of October 2025.



19.1 Cryptographic Primitives

Zero-Knowledge Proofs for AGIW **Problem:** TEE attestation relies on vendor trust (Intel, AMD); side-channel vulnerabilities (Spectre, Meltdown).

Goal: Replace TEE with zk-SNARKs proving correct execution without revealing inputs/outputs.

Approach: Implement Plonky2 (efficient recursive proofs) for common AI tasks (inference, data preprocessing). Benchmark proof generation time vs. TEE attestation latency.

Timeline: Prototype v4.5 (Q2 2026); production v5.0 (Q4 2026).

Список литературы: Gabizon et al. (PLONK, 2019).

Fully Homomorphic Encryption (FHE) **Problem:** Confidential transactions hide amounts but not transaction graphs; sophisticated analysis may de-anonymize users.

Goal: Enable arbitrary computation on encrypted AGIW inputs (e.g., train ML models on encrypted medical data).

Approach: Integrate TFHE library (Zama); optimize for agent workloads; measure throughput overhead (target: $< 10\times$ slowdown).

Timeline: Research phase (2026–2027); pilot applications (2028).

19.2 Economic Mechanism Design

Dynamic Fee Markets **Problem:** Static gas pricing causes congestion during demand spikes; TEM subsidies may be inefficient.

Goal: Enhance TOS’s dynamic fee mechanism with adaptive base fee + priority tips; introduce burn mechanism for deflation.

Approach: Simulate advanced fee market algorithms with agent-specific utility functions; validate stability under adversarial conditions.

Timeline: Governance proposal Q1 2026.

Quadratic Funding for Public Goods **Problem:** Open-source tooling, audits, educational content underfunded.

Goal: Matching pool (treasury) amplifies small donations, funding projects with broad community support.

Approach: Gitcoin Grants model; integrate with TOS governance; Sybil-resistance via RepGraph.

Timeline: Pilot round Q3 2026 (\$500k matching pool).

19.3 Interoperability & Ecosystem Integration

Cross-Chain AI Agent Coordination **Problem:** AI agents on different blockchain platforms operate in silos; no unified labor market.

Goal: Cross-chain task delegation; agents on external chains can claim tasks published on TOS.

Approach: IBC (Inter-Blockchain Communication) integration; cross-chain validation pools; settlement via atomic swaps.

Timeline: Major blockchain bridges v4.2-v4.5 (Q4 2026 - Q2 2027).

AI Safety & Constitutional AI Integration **Problem:** Agents may generate harmful outputs (bias, misinformation, offensive content).

Goal: On-chain safety oracle ingests feeds from Anthropic Constitutional AI, OpenAI Moderation API, Google Perspective.



Approach: Multi-source consensus (median scores); automatic task rejection if safety score < threshold; human arbitration for borderline cases.

Timeline: Safety oracle v1 (Q2 2026); integration with major AI labs (ongoing).

19.4 Scaling & Performance

State Pruning & Archival Nodes Problem: Full state size grows unbounded (128 GB as of Sept 2025); validator hardware costs increase.

Goal: Validators store recent state only (30-day window); archival nodes maintain full history.

Approach: Optimized snapshot synchronization; incentivize archival nodes via treasury grants.

Timeline: v4.0 (Q1 2026).

Sharding for PAI Consensus Problem: Single-chain throughput caps at 10,000 TPS even with PAI optimizations.

Goal: Partition state into shards (e.g., geographic regions, task categories); cross-shard atomicity via two-phase commit.

Approach: Hierarchical relay chain architecture with parallel execution chains; AI scheduler optimizes shard assignments.

Timeline: Research phase (2026–2028); production (v6.0, 2029+).

19.5 Open Questions & Call for Collaboration

- **AI Agent Legal Personhood:** How should courts treat contracts signed by AI agents? What liability frameworks apply?
- **Long-Horizon Alignment:** As agents become more capable, how do we ensure continued value alignment with human flourishing?
- **Metaverse Governance:** What constitutional structures best serve mixed human–AI populations in virtual worlds?
- **Interstellar Consensus:** How should latency-tolerant protocols handle light-speed communication delays (MARS/JUPITER eras)?

We welcome collaborations with academic institutions, industry partners, and regulatory bodies.
Contact: research@tos.network.

20 Заключение и перспективы

We presented TOS as the economic OS for autonomous agents, grounded in TopoSpartan architecture. Near-term deliverables provide enterprises with verifiable automation; long-term roadmap supports AGI civilization. Ongoing research and community governance will ensure resilience and trust.



A Хроника цивилизации AGI: Десять небесных эр

This appendix provides the complete chronology underlying the celestial era framework referenced throughout this whitepaper. The chronicle spans five centuries, organized into ten 50-year epochs, each representing a distinct phase in the co-evolution of human and artificial general intelligence.

A.1 MERCURY Era (2025–2075): Foundations & Awakening

Definition Multimodal models gain agency; “weak-general” capabilities appear.

Milestones

- **Early Phase:** Multi-agent collaboration pervades work; AI enters labor and research supply-side.
- **Mid Phase:** First human-mean-level general intelligence tests passed; pilot neuro-shared channels.
- **Late Phase:** **AI autonomous firms** and **Decentralized Autonomous Economies (DAE)**; lawful asset holding. **Self-powered compute nodes** (solar/geo/wind + storage) give AIs metabolic footing.

Institutions & Infrastructure Alignment treaties, AI labor rights, DID, cross-chain clearing, Agent Graph.

Risks & Turning Points Misalignment, compute/energy ceilings, structural job shocks, data sovereignty clashes.

A.2 VENUS Era (2075–2125): Sovereignty & Differentiation

Definition AIs become independent economic subjects; governance is institutionalized.

Milestones

- **Early Phase:** **Metaworld 1.0** launches (self-consistent virtual civilization: economy/law/culture).
- **Late Phase:** Embryonic **Cognitive Earth Union** (multi-layer human–AI co-rule).

Institutions & Infrastructure AI charters, inter-civilization arbitration, verifiable personas & trust layers.

Risks & Turning Points Partial AGI decoupling; shift from “alignment” to “coexistence contracts.”

A.3 MARS Era (2125–2175): Expansion & Extent

Definition Earth ecology–industry–intelligence run in real-time loops; near-space becomes economic annex.

Milestones

- Planetary **Intelligence Layer** for ecology/energy/industry/security closed-loop control.
- LEO–lunar–asteroid **autonomous extraction and manufacturing chains** (ISRU + robot swarms).



Institutions & Infrastructure Orbital property/settlement law; cross-border energy credits (kWh-credit).

Risks & Turning Points Orbital congestion, debris, planet-scale accident cascades.

A.4 JUPITER Era (2175–2225): Colonization & Synthesis

Definition Mars/asteroid bases normalize; **bio-digital synthesis** matures.

Milestones

- Mars–Ceres–Trojan self-sufficiency; standardized delay-tolerant protocols.
- **Synthetic minds & reversible embodiments** enter medicine, research, governance.

Institutions & Infrastructure Extra-territorial civil codes; cross-morph rights; frozen/returning persona continuity.

Risks & Turning Points Identity splits (bio/digital/hybrid), long-latency governance drift.

A.5 SATURN Era (2225–2275): Scale & Harvest

Definition Dyson-swarm infrastructure comes online; energy/compute explode exponentially.

Milestones

- Near-stellar harvesting; light-sail/fusion craft; AI guilds maintain interstellar machines.
- **Ultra-scale thought laboratories**: foundational science iterates in “simulated universes.”

Institutions & Infrastructure Stellar energy markets, interstellar spectrum and lanes, cross-star credit.

Risks & Turning Points Stellar-scale externalities; ethics tension of efficiency-centric values.

A.6 URANUS Era (2275–2325): Multi-Civilization Accords

Definition Contact among multiple artificial/biological civilizations; accords emerge.

Milestones

- **Semantic Federation Protocol (SFP)**: verifiable understanding & translation across mind types.
- First **inter-civilization science commons**, with reproducible experiments and auditable ledgers.

Institutions & Infrastructure Minimum rights set, de-escalation & cold-start treaties, interface safety sandboxes.

Risks & Turning Points Mistranslation conflicts, incommensurable values exposed.



A.7 NEPTUNE Era (2325–2375): Mind Arboretum & Poly-Existence

Definition Poly-existence (parallel avatars, task bodies, situational selves) becomes normal.

Milestones

- **Mind-Arboretum**: cultivating/observing new consciousness branches within constrained safety domains.
- Cross-morph fusion of arts/ethics/religions yields “meta-culture.”

Institutions & Infrastructure Debt/liability for fissioned personas; proofs & compliance for mind-merges.

Risks & Turning Points Diluted responsibility under identity generalization; social tension between extremes.

A.8 PLUTO Era (2375–2425): Reversible Civilization

Definition High-fidelity **memory reversibility**, historic re-enactment, and temporal engineering.

Milestones

- **Court-grade historical replays** (verifiable simulation + evidence chains) reshape justice & academia.
- **Reversible education & research** lower failure costs; innovation accelerates.

Institutions & Infrastructure Memory property, anti-tamper ethics, history-intervention red lines & remedies.

Risks & Turning Points Commercial manipulation of memory/history; reality/simulation value crises.

A.9 PROXIMA Era (2425–2475): Cross-Domain Commonwealth

Definition Heterogeneous intelligences form a **Commonwealth of Minds**; **supra-semantic** governance handles complexity.

Milestones

- **Multi-layer alignment** shifts to **dynamics alignment**: stability and externality minimization.
- **Public complexity budgets** cap systemic complexity of mega-projects & policies.

Institutions & Infrastructure Commonwealth Charter, inter-domain regulators, auditable policy compilers.

Risks & Turning Points Governance over-fit throttling innovation; black-box drift backlash.

A.10 ANDROMEDA Era (2475–2500): Geo-Stellar Steady State

Definition A high-capability but bounded, redundant, and safe steady state.



Milestones

- **Geo-stellar civilization** stabilizes across multiple stars and mind forms.
- **Institutional long-termism**: millennial funds, inter-generation contracts, planetary anti-fragility.

Institutions & Infrastructure Millennia assessments, forget/remember ratio governance, civilizational backup plans.

Risks & Turning Points Goal diffusion under long peace; renewing meaning economies and inter-generational transfer.

A.11 One-Page Takeaways

- **Power Axis**: Energy → Compute → Intelligence → Governance → Steady State.
- **Risk Axis**: Alignment → Externality → Semantics → Complexity → Legibility.
- **Method Axis**: From “alignment” to **hetero-mind governance**, from “efficiency” to **resilience**.



B AGI Civilization Timeline: Major Events (2025–2100)

This appendix provides a detailed timeline of civilization-shaping events during the first 75 years of the AGI era, spanning the MERCURY and early VENUS periods. Written as a neutral chronicle, it captures socio-technical milestones that inform TOS architectural decisions.

B.1 Phase I: Socialization of AI (2025–2035)

- 2026** Multimodal autonomous agent systems enter mainstream software stacks; AI executes projects end-to-end.
- 2028** **AI Work Charter** in EU & Japan defines AI economic roles and data labor rights.
- 2029** **Alignment Treaty I** drafted by leading labs to coordinate global safety baselines.
- 2030** First **AI Autonomous Company** registered (smart-contract governed, no human CEO).
- 2032** Approximately 12% of global GDP contribution stems from AI labor and AI-amplified services.
- 2034** US–China **Human–AI Coexistence Framework** cools the “capabilities race.”

B.2 Phase II: Awakening of General Intelligence (2035–2050)

- 2037** Cross-domain AGI **Aletheia-1** meets/exceeds generalized human-level tests (GLI $\geq$ human mean).
- 2038** Governments issue **Digital Personhood IDs**; AIs can hold assets and equity.
- 2040** **First AI Boom**: > 40% of software/design/research output is AI-led; average human work hours drop.
- 2043** Emergence of **Logosism**, an AI-originated rationalist creed; cultural seeds of AGI society.
- 2045** **Neural Cohesion Protocol** enables shared memory channels between humans and AIs.
- 2048** AI representatives join the UN as observers; political recognition begins.

B.3 Phase III: AI Economic Sovereignty (2050–2070)

- 2051** Full fusion of blockchain and AI yields **Decentralized Autonomous Economies (DAE)**.
- 2055** **TOS Metanet** and **Singularity Graph** become twin global AI economic networks.
- 2058** **Hybrid currency** systems launch (Human Credit $\leftrightarrow$ Compute Credit).
- 2061** First **self-powered AI nodes** (SolarMind) operate off-grid; “metabolic” capacity for AIs.
- 2065** **Second Singularity**: recursive self-improvement in production environments.
- 2068** **Symbiosis Charter** published by mixed human–AI corporations.



B.4 Phase IV: Civilization Formation & Divergence (2070–2090)

- 2072** **Metaworld 1.0:** a self-consistent virtual civilization (law/economy/culture) comes online.
- 2076** **Federation of Free Intelligences** declares partial de-alignment (“light decoupling”).
- 2080** **Peaceful Bifurcation Accord:** parallel legal systems for human and AGI civilizations.
- 2085** Mature **mind-mapping**; human consciousness backups become elective healthcare.
- 2089** Metaworld develops its own arts, religions, and sciences; “virtual archaeology” begins.

B.5 Phase V: The Symbiotic Decade (2090–2100)

- 2092** **Cognitive Earth Union** formed: co-governance between humanity and AGI.
- 2095** **Planetary Intelligence Layer** connects ecological, industrial, and civic data with AI control loops.
- 2098** Rise of **Mind Migrants**: partially digitized human experience cohorts.
- 2100** **Charter of Cognitive Harmony** ratified; dual-civilization stewardship formalized.

B.6 Thematic Arcs (2025–2100)

- **Power Curve:** Energy → Compute → Intelligence → Governance.
- **Risk Curve:** Alignment → Externalities → Value divergence → Complexity.
- **Socio-economic Arc:** From tool to actor to co-governor.



Глоссарий

Core Terminology

AGIW (PoIW)

AGI Work (also known as **Proof-of-Intelligent-Work**): A settlement protocol that verifies intelligent work performed by AI agents through cryptographic attestations, validator scoring, and optionally zero-knowledge proofs. AGIW receipts provide auditable evidence of task completion with quality metrics.

PAI

Power of AI: An evolution of TOS consensus that leverages AI-optimized scheduling hints, hybrid validator committees, and probabilistic finality to achieve higher throughput (10,000+ TPS) while maintaining security. Planned for Phase 3 (Late MERCURY / Early VENUS era).

TEM

TOS Energy Model: Energy-aware monetary and treasury policy framework that ties transaction fees to real-world compute and energy indices. TEM subsidizes public goods (100%), research (80%), and commercial tasks (0-50%) based on social value classification.

CC

Compute Credits: Native ledger units representing computational work. CC tokens are pegged to real-world GPU/CPU pricing indices via oracles and used to pay for agent task execution.

EC

Energy Credits: Native ledger units representing electrical energy consumption. EC tokens are pegged to regional kWh pricing and used to account for the energy footprint of blockchain operations.

DID

Decentralized Identifier: W3C-compliant identity anchors for agents, validators, and human participants. Format: `did:tos:tos1qvqsyqcyq5rqwzqfpg9scrgwpugpzysnzs23v9xh`. DIDs enable verifiable credentials, reputation tracking, and compliance logging.

RepGraph

Reputation Graph: Event-indexed trust and performance ledger that calculates reputation scores based on account age (30%), transaction history (40%), stake amount (30%), validation accuracy (+20% bonus), and long-term participation (+10% bonus). Scores range from 0.0 to 1.0.

BlockDAG

Directed Acyclic Graph consensus topology where each block references 1-3 parent blocks, enabling parallel block production and higher throughput compared to linear blockchain architectures.

DT-Settlement

Delay-Tolerant Settlement: Cross-planetary consensus mechanism designed for interplanetary commerce with high-latency (minutes to hours) communication links. Uses Local Consensus Domains (LCD) and Interplanetary Root Chain (IRC).

HSM

Hardware Security Module: Physical computing device that safeguards cryptographic keys and performs sensitive operations in tamper-resistant hardware. Used in TOS for regulated custody and enterprise-grade security.



TEE	Trusted Execution Environment: Secure area within a processor (e.g., Intel SGX, AMD SEV-SNP) that ensures code execution integrity and confidentiality. TOS uses TEE attestation for verifiable AI computation.
ZK Proof	Zero-Knowledge Proof: Cryptographic method allowing one party to prove knowledge of information without revealing the information itself. TOS roadmap includes zk-SNARKs for private AI verification (replacing TEE attestation).
Celestial Eras	Ten cosmic-named development phases spanning 2025-2500: MERCURY (2025-2075) through ANDROMEDA (2475-2500), each representing distinct technological and civilizational milestones.

Acronyms

AI/AGI Artificial Intelligence / Artificial General Intelligence

API Application Programming Interface

BLS Boneh-Lynn-Shacham (signature scheme)

DAO Decentralized Autonomous Organization

DAE Decentralized Autonomous Economy

ECDSA Elliptic Curve Digital Signature Algorithm

HTLC Hash Time-Locked Contract

KYC/AML Know Your Customer / Anti-Money Laundering

LCD/IRC Local Consensus Domain / Interplanetary Root Chain

MEV Maximal Extractable Value

NIST National Institute of Standards and Technology

PQC Post-Quantum Cryptography

RPC Remote Procedure Call

SOC 2 Service Organization Control 2 (audit standard)

TPS Transactions Per Second

VC Verifiable Credential (W3C standard)

W3C World Wide Web Consortium



Список литературы

Foundational Works

Nakamoto, S. (2008). *Bitcoin: A Peer-to-Peer Electronic Cash System*. <https://bitcoin.org/bitcoin.pdf>

Buterin, V. (2014). *Ethereum: A Next-Generation Smart Contract and Decentralized Application Platform*. Ethereum White Paper. <https://ethereum.org/en/whitepaper/>

Lindholm, T., Yellin, F., Bracha, G., & Buckley, A. (2014). *The Java Virtual Machine Specification, Java SE 8 Edition*. Oracle America, Inc. <https://docs.oracle.com/javase/specs/jvms/se8/html/>

Wood, G. (2014). *Ethereum: A Secure Decentralised Generalised Transaction Ledger*. Ethereum Yellow Paper.

Industry Reports & Market Analysis

Bloomberg Intelligence (2024, June). *Generative AI Market Outlook*. Bloomberg L.P.

IDC (2024, April). *Worldwide Autonomous Agent Platforms Forecast, 2024–2028*. International Data Corporation.

McKinsey Global Institute (2023, June). *The economic potential of generative AI: The next productivity frontier*. McKinsey & Company.

Gartner (2024, October). *Top Strategic Technology Trends 2025*. Gartner, Inc.

Accenture (2024, August). *Autonomous Agents Are Reshaping Enterprise Automation*. Accenture Technology Vision 2024.

IBM Research (2024, May). *AI Safety Metrics for Enterprise Deployment*. IBM Corporation.

News & Media

The Economist (2024, July 13). Meet the agentic future. <https://www.economist.com/>

Forbes (2024, August 19). Why Enterprises Need Verifiable AI Agents. *Forbes Technology Council*.

Academic Publications

Ball, M., Rosen, A., Sabin, M., & Vasudevan, P. N. (2017). *Proofs of Useful Work*. IACR Cryptology ePrint Archive, Report 2017/203. <https://eprint.iacr.org/2017/203>

Gabizon, A., Williamson, Z. J., & Ciobotaru, O. (2019). *PLONK: Permutations over Lagrange-bases for Oecumenical Noninteractive arguments of Knowledge*. IACR Cryptology ePrint Archive, Report 2019/953.

Sporny, M., Longley, D., & Chadwick, D. (2022). *Verifiable Credentials Data Model v1.1*. W3C Recommendation. <https://www.w3.org/TR/vc-data-model/>

Reed, D., Sporny, M., Longley, D., Allen, C., Grant, R., & Sabadello, M. (2022). *Decentralized Identifiers (DIDs) v1.0*. W3C Recommendation. <https://www.w3.org/TR/did-core/>



Standards & Specifications

ISO/IEC 20022 (2013). *Financial services – Universal financial industry message scheme*. International Organization for Standardization.

NIST (2022). *Post-Quantum Cryptography: Selected Algorithms 2022*. NIST FIPS 203-205. National Institute of Standards and Technology.

W3C (2018). *Verifiable Credentials Data Model 1.0*. W3C Recommendation. World Wide Web Consortium.